



See what your computer is doing

With Ftrace utilities

Steven Rostedt

Open Source Engineer

rostedt@goodmis.org / srostedt@vmware.com

Twitter: @VMWopensource / Expo floor #31

This is a tutorial

Please follow along to get the most out of it

Most distributions have Ftrace configured in their kernels

- I'm using the default Debian “testing” kernel
(4.18.0-3-amd64)

Must be “root”

- No “sudo” for now

Helpful utilities for later:

- trace-cmd
- KernelShark

git clone /media/<user>/FTRACETUT/trace-cmd.git

Where do we find ftrace?

Most distributions mount: `/sys/kernel/debug`

- find the “tracing” directory there

Can also mount the tracing directory directly

- `mount -t tracefs nodev /sys/kernel/tracing`

We'll use `/sys/kernel/tracing` here

Let's look at the tracing directory

```
# mount -t tracefs nodev /sys/kernel/tracing
# cd /sys/kernel/tracing
# ls

available_events      per_cpu              stack_trace
available_filter_functions  printk_formats      stack_trace_filter
available_tracers    README             timestamp_mode
buffer_size_kb        saved_cmdlines       trace
buffer_total_size_kb  saved_cmdlines_size  trace_clock
current_tracer      saved_tgids          trace_marker
dyn_ftrace_total_info set_event             trace_marker_raw
enabled_functions     set_event_pid        trace_options
events              set_ftrace_filter  trace_pipe
free_buffer           set_ftrace_notrace   tracing_cpumask
instances             set_ftrace_pid       tracing_max_latency
kprobe_events         set_graph_function tracing_on
kprobe_profile        set_graph_notrace    tracing_thresh
max_graph_depth       snapshot             uprobe_events
options               stack_max_size       uprobe_profile
```

README

Actually lets you know if some things are available

- Not everything, but features are only in here if they exist in the kernel

Has basic reference on how to perform some features

Has basic information to some of the files in tracefs

available_tracers

A “tracer” performs some action

- nop - No action (no tracer enabled)
- function - Trace kernel functions
- function_graph - Graph kernel functions
- blk - Used with blktrace (not discussed in this tutorial)
- mmiotrace - Trace interactions between drivers and hardware (not discussed in this tutorial)

Other available_tracers

These tracers are not available in my Debian kernel

- hwlat - Detects hardware latency (i.e. SMI interrupts)
- wakeup{,_dl,_rt} - Records max wake up latency
- irqsoff - Records max interrupt latency
- preemptoff - Records max preemption disabled latency
- preemptirqsoff - Records max preemption or interrupt disabled latency

We wont talk about these in this tutorial either

- But they are good to know about

current_tracer

The way to enable a tracer

```
# echo function > current_tracer  
# echo nop > current_tracer
```


trace

Where to see the output of the trace

```
# echo function > current_tracer
# cat trace

# tracer: function
#
#          _-----=> irqsoft
#          / _-----=> need-resched
#          | / _-----=> hardirq/softirq
#          || / _-----=> preempt-depth
#          ||| / _-----=> delay
#          |||| /
#
# TASK-PID   CPU#   |         |   TIMESTAMP   FUNCTION
#             |   |   |         |   |
soffice.bin-1850 [003] .... 1277817.313622: __fget_light <-sockfd_lookup_light
soffice.bin-1850 [003] .... 1277817.313622: __fget <-__fget_light
soffice.bin-1850 [003] .... 1277817.313623: __sys_recvmsg <-__sys_recvmsg
soffice.bin-1850 [003] .... 1277817.313623: copy_msghdr_from_user <-__sys_recvmsg
soffice.bin-1850 [003] .... 1277817.313624: rw_copy_check_uvector <-import_iovec
soffice.bin-1850 [003] .... 1277817.313624: __check_object_size <-rw_copy_check_uvector
soffice.bin-1850 [003] .... 1277817.313624: __virt_addr_valid <-__check_object_size
soffice.bin-1850 [003] .... 1277817.313625: check_stack_object <-__check_object_size
soffice.bin-1850 [003] .... 1277817.313625: sock_recvmsg <-__sys_recvmsg
soffice.bin-1850 [003] .... 1277817.313625: security_socket_recvmsg <-sock_recvmsg
soffice.bin-1850 [003] .... 1277817.313626: apparmor_socket_recvmsg <-security_socket_recvmsg
```

trace

Remember to turn off the tracer

```
# echo nop > current_tracer
# cat trace

# tracer: nop
#
#
#          _-----=> irqsoff
#         / _-----=> need_resched
#        | / _-----=> hardirq/softirq
#       || / _-----=> preempt-depth
#      ||| / _-----=> delay
#     |||| / _-----=>
#
# TASK-PID   CPU#  | TIMESTAMP | FUNCTION |
#   | |       |   |          |           |
```

tracing_on

Stop updates to the trace without clearing it.

Note, tracing is still happening, just not recording (will have overhead)

```
# echo function > current_tracer
# echo 0 > tracing_on
# cat trace

# tracer: function
#
#          _-----=> irqsoff
#          / _-----=> need-resched
#          | / _----=> hardirq/softirq
#          || / _--=> preempt-depth
#          ||| /      delay
#          |||| /
#          TASK-PID  CPU#  | |||   TIMESTAMP  FUNCTION
#          | |       |   | |||   |
xfwm4-17105 [001] d... 1278457.346868: hrtick_update <-ttwu_do_activate
xfwm4-17105 [001] d... 1278457.346868: ttwu_do_wakeup <-try_to_wake_up
xfwm4-17105 [001] d... 1278457.346869: check_preempt_curr <-ttwu_do_wakeup
xfwm4-17105 [001] d... 1278457.346869: resched_curr <-check_preempt_curr
xfwm4-17105 [001] d... 1278457.346870: ttwu_stat <-try_to_wake_up
xfwm4-17105 [001] d... 1278457.346870: _raw_spin_unlock_irqrestore <-try_to_wake_up
xfwm4-17105 [001] d... 1278457.346871: _raw_spin_unlock_irqrestore <-ep_poll_callback
xfwm4-17105 [001] d... 1278457.346871: _raw_spin_unlock_irqrestore <-__wake_up_common_lock
```

tracing_on (Caution)

Writing ASCII “0” into the file, stops tracing

Writing ASCII “1” into the file starts tracing

Note, be sure to add a space between the number and ‘>’

Do this:

```
# echo 0 > tracing_on  
# echo 1 > tracing_on
```

Not this:

```
# echo 0>tracing_on  
# echo 1>tracing_on
```

tracing_on (Caution)

Writing ASCII “0” into the file, stops tracing

Writing ASCII “1” into the file starts tracing

Note, be sure to add a space between the number and ‘>’

Do this:

```
# echo 0 > tracing_on  
# echo 1 > tracing_on
```

Not this:

```
# echo 0>tracing_on  
# echo 1>tracing_on
```

Why?

tracing_on (Caution)

```
# echo 0>tracing_on  
# echo 1>tracing_on
```

Writes standard input and standard output into tracing_on respectively

trace_marker

Writes into the ftrace ring buffer (along with other kernel events)

```
# echo nop > current_tracer
# echo 1 > tracing_on
# echo 'hello world!' > trace_marker
# cat trace

# tracer: nop
#
#          _-----=> irqsoff
#         / _-----=> need_resched
#        | / _-----=> hardirq/softirq
#       || / _-----=> preempt-depth
#      ||| / _-----=> delay
#     TASK-PID  CPU#  | |||   TIMESTAMP  FUNCTION
#     | |       |   | |||   |          |
bash-13441 [003] .... 1279231.118528: tracing_mark_write: hello world!
```

trace_marker

In C code:

```
static int marker_fd = -1;

static void setup_trace_marker(void)
{
    marker_fd = open("/sys/kernel/tracing/trace_marker", O_WRONLY);
}

static void write_trace_marker(const char *fmt, ...)
{
    char buf[BUFSIZ];
    va_list ap;
    int n;

    if (marker_fd < 0)
        return;

    va_start(ap, fmt);
    n = vsnprintf(buf, BUFSIZ, fmt, ap);
    va_end(ap);

    write(marker_fd, buf, n);
}
```


Function tracing

Trace almost any function in the kernel!

- See the possible functions in “available_filter_functions”

Limit what functions you trace

- set_ftrace_filter - Only trace these functions
- set_ftrace_notrace - Do not trace these functions
(**notrace** takes precedence over **filter**)

set_fttrace_filter and set_fttrace_notrace

Just echo function names into the files

```
# echo foo > set_fttrace_notrace
```

Can add more than one at a time (white space delimited)

```
# echo foo bar > set_fttrace_filter
```

Append with the bash concatenation ">>"

```
# echo zoot >> set_fttrace_filter
```

Clear with just writing nothing into it

```
# echo > set_fttrace_notrace
```

set_fttrace_filter and set_fttrace_notrace

Can handle minor wild cards “*” and “?”

```
# echo '?lock*d' > set_fttrace_filter
```

Can use available_filter_functions for more complex filtering

```
# cut -d' ' -f1 available_filter_functions | grep -E 'ipv(4|6)' > set_fttrace_filter
```

- Note, the ‘cut’ is to remove module names:

```
nf_reject_ip_tcphdr_get [nf_reject_ipv4]
```

Function tracing

```
# echo '*lock*' > set_ftrace_filter
# echo '*clock*' > set_ftrace_notrace
# echo function > current_tracer
# cat trace

# tracer: function
#
#          _-----=> irqs-off
#          / _-----=> need-resched
#          | / _---=> hardirq/softirq
#          || / _--=> preempt-depth
#          ||| /      delay
#          ||||
# TASK-PID  CPU#  | TIMESTAMP | FUNCTION
#   |   |   |   |   |   |
Xorg-16967 [002] .... 1383242.709737: mutex_trylock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709737: ww_mutex_unlock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709737: mutex_unlock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709738: _raw_spin_lock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709738: mutex_trylock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709738: ww_mutex_unlock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709739: mutex_unlock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709739: _raw_spin_lock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709740: mutex_trylock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709740: ww_mutex_unlock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709740: mutex_unlock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709741: _raw_spin_lock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709741: mutex_trylock <-i915_vma_retire
Xorg-16967 [002] .... 1383242.709741: ww_mutex_unlock <-i915_vma_retire
```

Function Graph tracing

Similar to function tracing (traces the same functions)

Also traces the return of a function

Allows to see a graph of the function calls

- What function called which function

Function Graph tracing

```
# echo function_graph > current_tracer
# cat trace

# tracer: function_graph
#
# CPU    DURATION    FUNCTION CALLS
# |      |      |      |      |      |
2) 0.032 us |      } /* fput */
2) 7.984 us |      } /* __sys_recvmsg */
2) 8.259 us |      } /* __x64_sys_recvmsg */
2) 8.558 us |      } /* do_syscall_64 */
2)          | do_syscall_64() {
2)          |     __x64_sys_poll() {
2)          |         poll_select_set_timeout() {
2) 0.104 us |             ktime_get_ts64();
2) 0.045 us |             timespec64_add_safe();
2) 0.738 us |         }
2)          |     do_sys_poll() {
2)          |         __check_object_size() {
2) 0.034 us |             __virt_addr_valid();
2) 0.036 us |             check_stack_object();
2) 0.613 us |         }
2)          |     select_estimate_accuracy() {
2) 0.053 us |         ktime_get_ts64();
2) 0.033 us |         set_normalized_timespec64();
2) 0.681 us |     }
2)          |     __fdget() {
2)          |         __fget_light() {
2) 0.069 us |             __fget();
```

Function Graph tracing

```
# echo function_graph > current_tracer
# cat trace

# tracer: function_graph
#
# CPU    DURATION    FUNCTION CALLS
# |      |      |      |      |      |
2) 0.032 us      } /* fput */
2) 7.984 us      } /* __sys_recvmsg */
2) 8.259 us      } /* __x64_sys_recvmsg */
2) 8.558 us      } /* do_syscall_64 */
2) do_syscall_64() {
2)   __x64_sys_poll() {
2)     poll_select_set_timeout() {
2)       0.104 us   ktime_get_ts64();
2)       0.045 us   timespec64_add_safe();
2)       0.738 us   }
2)     do_sys_poll() {
2)       __check_object_size() {
2)         0.034 us   __virt_addr_valid();
2)         0.036 us   check_stack_object();
2)         0.613 us   }
2)     select_estimate_accuracy() {
2)       0.053 us   ktime_get_ts64();
2)       0.033 us   set_normalized_timespec64();
2)       0.681 us   }
2)     __fdget() {
2)       __fget_light() {
2)         0.069 us   __fget();
```

Function Graph tracing

Don't trust all the times (just take it as a guide)

Function graph tracer adds overhead

- You see the overhead of functions traced within other functions

Closest to actual time is a single entity

- The enter and exit of a function are together
- Denoted by single event with ';'
- Instead of two events with '{' and '}'

The `set_ftrace_filter` and `set_ftrace_notrace` affect `function_graph`

Function Graph tracing

```
# echo do_IRQ > set_ftrace_filter
# echo function_graph > current_tracer
# cat trace

# tracer: function_graph
#
# CPU    DURATION                FUNCTION CALLS
# |      |      |                |      |      |      |
2) =====> |
2) + 14.098 us | do_IRQ();
2) <===== |
2) =====> |
2) 6.920 us  | do_IRQ();
2) <===== |
2) =====> |
2) 6.259 us  | do_IRQ();
2) <===== |
2) =====> |
2) 5.625 us  | do_IRQ();
2) <===== |
2) =====> |
2) + 10.433 us | do_IRQ();
2) <===== |
2) =====> |
2) + 11.347 us | do_IRQ();
2) <===== |
2) =====> |
2) + 11.230 us | do_IRQ();
```

Tracing Events

Function tracing is great BUT!

- It doesn't show much data besides the functions being called
- No parameters or variables can be seen

Which brings us to Trace Events

- They are points in the kernel that write into the trace buffer
- Record specific data within the kernel
- Allows to see more detailed view of what is happening

Tracing Events

They are grouped by “system”

- sched - schedule events
- irq - interrupt events
- net - networking events
- syscalls - system call events (all system calls and their parameters)
- module - module loading, unloading, freeing, etc
- kvm - events for the KVM guests
- exceptions - page faults
- cgroup - changes to cgroups
- And many many more!

Trace Event Systems

```
# ls events
```

alarmtimer	fib6	iommu	napi	rseq	udp
asoc	filelock	irq	net	rtc	v4l2
block	filemap	irq_matrix	nmi	sched	vb2
bridge	fs_dax	irq_vectors	oom	scsi	vmscan
btrfs	ftrace	jbd2	page_isolation	signal	vsyscall
cfg80211	gpio	kmem	pagemap	skb	wbt
cgroup	hda	kvm	percpu	smbus	workqueue
clk	hda_controller	kvmmmu	power	sock	writeback
compaction	hda_intel	libata	printk	spi	x86_fpu
cpuhp	header_event	mac80211	qdisc	sunrpc	xdp
devlink	header_page	mce	random	swiotlb	xen
dma_fence	huge_memory	mei	ras	syscalls	xfs
drm	hyperv	migrate	raw_syscalls	task	xhci-hcd
enable	i2c	mmc	rcu	tcp	
exceptions	i915	module	regmap	thermal	
ext4	initcall	mpx	regulator	timer	
fib	intel-sst	msr	rpm	tlb	

Tracing Events

Each system has several events

For example: The “sched” system

- sched_switch
- sched_waking
- sched_process_fork
- sched_process_exit
- sched_process_exec
- sched_migrate_task
- etc

Trace Events

```
# ls events/sched
```

enable	sched_process_fork	sched_stick_numa
filter	sched_process_free	sched_swap_numa
sched_kthread_stop	sched_process_hang	sched_switch
sched_kthread_stop_ret	sched_process_wait	sched_wait_task
sched_migrate_task	sched_stat_blocked	sched_wake_idle_without_ipi
sched_move_numa	sched_stat_iowait	sched_wakeup
sched_pi_setprio	sched_stat_runtime	sched_wakeup_new
sched_process_exec	sched_stat_sleep	sched_waking
sched_process_exit	sched_stat_wait	

Trace Events

```
# ls events/sched
```

```
enable
```

```
filter
```

```
sched_kthread_stop
```

```
sched_kthread_stop_ret
```

```
sched_migrate_task
```

```
sched_move_numa
```

```
sched_pi_setprio
```

```
sched_process_exec
```

```
sched_process_exit
```

```
sched_process_fork
```

```
sched_process_free
```

```
sched_process_hang
```

```
sched_process_wait
```

```
sched_stat_blocked
```

```
sched_stat_iowait
```

```
sched_stat_runtime
```

```
sched_stat_sleep
```

```
sched_stat_wait
```

```
sched_stick_numa
```

```
sched_swap_numa
```

```
sched_switch
```

```
sched_wait_task
```

```
sched_wake_idle_without_ipi
```

```
sched_wakeup
```

```
sched_wakeup_new
```

```
sched_waking
```

Enabling Trace Events

```
# echo nop > current_tracer
# echo 1 > events/sched/sched_waking/enable
# echo 1 > events/sched/sched_switch/enable
# cat trace
```

```
# tracer: nop
#
#
#          _-----> irqs-off
#          / _-----> need-resched
#          / _-----> hardirq/softirq
#          || / _-----> preempt-depth
#          ||| / _-----> delay
#          ||| /
#
# TASK-PID   CPU#  | TIMESTAMP | FUNCTION
# |-----|-----|-----|-----|
Timer-19729 [000] d.h. 1555860.848766: sched_waking: comm=ModuleProcessTh pid=24498 prio=120 target_cpu=000
Timer-19729 [000] d... 1555860.848777: sched_waking: comm=Web Content pid=19688 prio=120 target_cpu=002
Timer-19729 [000] d... 1555860.848818: sched_switch: prev_comm=Timer prev_pid=19729 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_prio=120
<idle>-0 [000] d... 1555860.848862: sched_switch: prev_comm=swapper/0 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=Xorg next_pid=16967 next_prio=120
Xorg-16967 [000] d... 1555860.848917: sched_switch: prev_comm=Xorg prev_pid=16967 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_prio=120
<idle>-0 [000] d.h. 1555860.850480: sched_waking: comm=SoftwareVsyncTh pid=19636 prio=120 target_cpu=000
<idle>-0 [000] d... 1555860.850513: sched_switch: prev_comm=swapper/0 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=SoftwareVsyncTh next_pid=19636 next_prio=
SoftwareVsyncTh-19636 [000] d... 1555860.850562: sched_waking: comm=IPDL Background pid=19629 prio=120 target_cpu=002
SoftwareVsyncTh-19636 [000] d... 1555860.850604: sched_switch: prev_comm=SoftwareVsyncTh prev_pid=19636 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_prio=
<idle>-0 [000] d... 1555860.850629: sched_switch: prev_comm=swapper/0 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=Gecko_IOThread next_pid=19602 next_prio=
Gecko_IOThread-19602 [000] d... 1555860.850677: sched_waking: comm=Chrome_ChildThr pid=19690 prio=120 target_cpu=001
Gecko_IOThread-19602 [000] d... 1555860.850701: sched_switch: prev_comm=Gecko_IOThread prev_pid=19602 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_prio=
<idle>-0 [000] d.h. 1555860.861078: sched_waking: comm=InputThread pid=16979 prio=120 target_cpu=001
<idle>-0 [000] d... 1555860.861311: sched_switch: prev_comm=swapper/0 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=Xorg next_pid=16967 next_prio=120
Xorg-16967 [000] d... 1555860.861384: sched_switch: prev_comm=Xorg prev_pid=16967 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_prio=120
<idle>-0 [000] d.h. 1555860.863389: sched_waking: comm=SendControllerT pid=5426 prio=120 target_cpu=000
<idle>-0 [000] d... 1555860.863418: sched_switch: prev_comm=swapper/0 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=SendControllerT next_pid=5426 next_prio=
SendControllerT-5426 [000] d... 1555860.863480: sched_switch: prev_comm=SendControllerT prev_pid=5426 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_prio=
<idle>-0 [000] d.h. 1555860.863524: sched_waking: comm=SendControllerT pid=5426 prio=120 target_cpu=000
<idle>-0 [000] d... 1555860.863535: sched_switch: prev_comm=swapper/0 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=SendControllerT next_pid=5426 next_prio=
SendControllerT-5426 [000] d... 1555860.863571: sched_waking: comm=TaskSchedulerFo pid=27478 prio=120 target_cpu=001
```


Enabling Trace Events

```
sched_waking: comm=ModuleProcessTh pid=24498 prio=120 target_cpu=000
sched_waking: comm=Web Content pid=19688 prio=120 target_cpu=002
sched_switch: prev_comm=Timer prev_pid=19729 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_prio=120
sched_switch: prev_comm=swapper/0 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=Xorg next_pid=16967 next_prio=120
sched_switch: prev_comm=Xorg prev_pid=16967 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_prio=120
sched_waking: comm=SoftwareVsyncTh pid=19636 prio=120 target_cpu=000
sched_switch: prev_comm=swapper/0 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=SoftwareVsyncTh next_pid=19636 next_p
sched_waking: comm=IPDL Background pid=19629 prio=120 target_cpu=002
sched_switch: prev_comm=SoftwareVsyncTh prev_pid=19636 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_p
sched_switch: prev_comm=swapper/0 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=Gecko_IOThread next_pid=19602 next_pr
sched_waking: comm=Chrome_ChildThr pid=19690 prio=120 target_cpu=001
sched_switch: prev_comm=Gecko_IOThread prev_pid=19602 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_pr
sched_waking: comm=InputThread pid=16979 prio=120 target_cpu=001
sched_switch: prev_comm=swapper/0 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=Xorg next_pid=16967 next_prio=120
sched_switch: prev_comm=Xorg prev_pid=16967 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_prio=120
sched_waking: comm=SendControllerT pid=5426 prio=120 target_cpu=000
sched_switch: prev_comm=swapper/0 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=SendControllerT next_pid=5426 next_pr
sched_switch: prev_comm=SendControllerT prev_pid=5426 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_pr
sched_waking: comm=SendControllerT pid=5426 prio=120 target_cpu=000
sched_switch: prev_comm=swapper/0 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=SendControllerT next_pid=5426 next_pr
sched_waking: comm=TaskSchedulerFo pid=27478 prio=120 target_cpu=001
```

Enabling Groups of Events

```
# echo nop > current_tracer
# echo 1 > events/sched/enable
# cat trace
```

```
# tracer: nop
```

```
# #          /-----> irqsoft-off
# #          /-----> need-resched
# #          | /-----> hardirq/softirq
# #          || /-----> preempt-depth
# #          ||| /-----> delay
# #          TASK-PID   CPU#      TIMESTAMP    FUNCTION
# #          | |         |         |           |
firefox.real-19591 [002] d... 1558025.460824: sched_stat_runtime: comm=JS Helper pid=19608 runtime=7550 [ns] vruntime=111608395027855 [ns]
firefox.real-19591 [002] d... 1558025.460826: sched_wakeup: comm=JS Helper pid=19607 prio=120 target_cpu=000
firefox.real-19591 [002] d... 1558025.460827: sched_waking: comm=JS Helper pid=19606 prio=120 target_cpu=000
firefox.real-19591 [002] d... 1558025.460829: sched_stat_runtime: comm=JS Helper pid=19607 runtime=2756 [ns] vruntime=111608395161002 [ns]
firefox.real-19591 [002] d... 1558025.460830: sched_wakeup: comm=JS Helper pid=19606 prio=120 target_cpu=000
firefox.real-19591 [002] d... 1558025.460831: sched_waking: comm=JS Helper pid=19611 prio=120 target_cpu=000
firefox.real-19591 [002] d... 1558025.460833: sched_stat_runtime: comm=JS Helper pid=19606 runtime=2174 [ns] vruntime=111608394990889 [ns]
firefox.real-19591 [002] d... 1558025.460834: sched_wakeup: comm=JS Helper pid=19611 prio=120 target_cpu=000
firefox.real-19591 [002] d... 1558025.460835: sched_waking: comm=JS Helper pid=19609 prio=120 target_cpu=000
firefox.real-19591 [002] d... 1558025.460837: sched_stat_runtime: comm=JS Helper pid=19611 runtime=2218 [ns] vruntime=111608394992422 [ns]
firefox.real-19591 [002] d... 1558025.460838: sched_wakeup: comm=JS Helper pid=19609 prio=120 target_cpu=000
firefox.real-19591 [002] d... 1558025.460839: sched_waking: comm=JS Helper pid=19608 prio=120 target_cpu=000
firefox.real-19591 [002] d... 1558025.460841: sched_stat_runtime: comm=JS Helper pid=19609 runtime=2484 [ns] vruntime=111608394997215 [ns]
firefox.real-19591 [002] d... 1558025.460841: sched_wakeup: comm=JS Helper pid=19608 prio=120 target_cpu=000
firefox.real-19591 [002] d... 1558025.461027: sched_waking: comm=JS Helper pid=19608 prio=120 target_cpu=000
firefox.real-19591 [002] d... 1558025.461029: sched_wake_idle_without_ipi: cpu=0
firefox.real-19591 [002] d... 1558025.461030: sched_wakeup: comm=JS Helper pid=19608 prio=120 target_cpu=000
firefox.real-19591 [002] d... 1558025.461030: sched_waking: comm=JS Helper pid=19606 prio=120 target_cpu=001
firefox.real-19591 [002] d... 1558025.461032: sched_wake_idle_without_ipi: cpu=1
firefox.real-19591 [002] d... 1558025.461032: sched_wakeup: comm=JS Helper pid=19606 prio=120 target_cpu=001
firefox.real-19591 [002] d... 1558025.461032: sched_waking: comm=JS Helper pid=19611 prio=120 target_cpu=000
```

Enabling All Events!

```
# echo 1 > events/enable
# cat trace

# tracer: nop
#
#
#          _-----> irqsoft
#          /_-----> need-resched
#          | /_-----> hardirq/softirq
#          || /_-----> preempt-depth
#          ||| /_-----> delay
#
# TASK-PID   CPU#  TIMESTAMP  FUNCTION
#   |   |   |   |   |   |
hostd-fair-2334 [003] d... 1558119.421939: tlb_flush: pages:0 reason:flush on task switch (0)
hostd-fair-2334 [003] d... 1558119.421940: x86_fpu_regs_deactivated: x86/fpu: 000000000047327aa initialized: 1 xfeatures: 3 xcomp_bv: 800000000000001f
<idle>-0 [003] d... 1558119.421942: cpu_idle: state=2 cpu_id=3
<idle>-0 [003] dN.. 1558119.422000: cpu_idle: state=4294967295 cpu_id=3
<idle>-0 [003] dN.. 1558119.422001: rcu_utilization: Start context switch
<idle>-0 [003] dN.. 1558119.422001: rcu_utilization: End context switch
<idle>-0 [003] d... 1558119.422003: sched_switch: prev_comm=swapper/3 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=hostd-fair next_pid=
<idle>-0 [003] d... 1558119.422003: tlb_flush: pages:0 reason:flush on task switch (0)
<idle>-0 [003] d... 1558119.422004: write_msr: c0000100, value 7f8741977700
<idle>-0 [003] d... 1558119.422004: x86_fpu_regs_activated: x86/fpu: 000000000047327aa initialized: 1 xfeatures: 3 xcomp_bv: 800000000000001f
hostd-fair-2334 [003] .... 1558119.422005: sys_exit: NR 202 = 0
hostd-fair-2334 [003] .... 1558119.422005: sys_futex -> 0x0
hostd-fair-2334 [003] .... 1558119.422007: sys_enter: NR 202 (5558d741eae0, 81, 1, 0, 0, 5558d6f7bcd8)
hostd-fair-2334 [003] .... 1558119.422007: sys_futex(uaddr: 5558d741eae0, op: 81, val: 1, utime: 0, uaddr2: 0, val3: 5558d6f7bcd8)
hostd-fair-2334 [003] .... 1558119.422008: sys_exit: NR 202 = 0
hostd-fair-2334 [003] .... 1558119.422008: sys_futex -> 0x0
hostd-fair-2334 [003] .... 1558119.422010: sys_enter: NR 202 (5558d741eb3c, 80, 0, 0, 0, 5558d6f7bcd8)
hostd-fair-2334 [003] .... 1558119.422011: sys_futex(uaddr: 5558d741eb3c, op: 80, val: 0, utime: 0, uaddr2: 0, val3: 5558d6f7bcd8)
hostd-fair-2334 [003] d... 1558119.422011: rcu_utilization: Start context switch
hostd-fair-2334 [003] d... 1558119.422012: rcu_utilization: End context switch
hostd-fair-2334 [003] d... 1558119.422012: sched_stat_runtime: comm=hostd-fair pid=2334 runtime=13739 [ns] vruntime=80967984047 [ns]
```

What's in an event?

```
# cat events/sched/sched_switch/format
```

```
name: sched_switch
```

```
ID: 312
```

```
format:
```

```
    field:unsigned short common_type;    offset:0;    size:2;    signed:0;
    field:unsigned char common_flags;    offset:2;    size:1;    signed:0;
    field:unsigned char common_preempt_count; offset:3;    size:1;    signed:0;
    field:int common_pid;    offset:4;    size:4;    signed:1;
```

```
    field:char prev_comm[16];    offset:8;    size:16;    signed:1;
    field:pid_t prev_pid;    offset:24;    size:4;    signed:1;
    field:int prev_prio;    offset:28;    size:4;    signed:1;
    field:long prev_state; offset:32;    size:8;    signed:1;
    field:char next_comm[16];    offset:40;    size:16;    signed:1;
    field:pid_t next_pid;    offset:56;    size:4;    signed:1;
    field:int next_prio;    offset:60;    size:4;    signed:1;
```

```
print fmt: "prev_comm=%s prev_pid=%d prev_prio=%d prev_state=%s%s ==> next_comm=%s next_pid=%d
next_prio=%d", REC->prev_comm, REC->prev_pid, REC->prev_prio, (REC->prev_state & (((0x0000 | 0x0001 |
0x0002 | 0x0004 | 0x0008 | 0x0010 | 0x0020 | 0x0040) + 1) << 1) - 1) ? __print_flags(REC->prev_state
& (((0x0000 | 0x0001 | 0x0002 | 0x0004 | 0x0008 | 0x0010 | 0x0020 | 0x0040) + 1) << 1) - 1), "|",
{ 0x01, "S" }, { 0x02, "D" }, { 0x04, "T" }, { 0x08, "t" }, { 0x10, "X" }, { 0x20, "Z" }, { 0x40,
"P" }, { 0x80, "I" }) : "R", REC->prev_state & (((0x0000 | 0x0001 | 0x0002 | 0x0004 | 0x0008 | 0x0010
| 0x0020 | 0x0040) + 1) << 1) ? "+" : "", REC->next_comm, REC->next_pid, REC->next_prio
```

What's in an event?

```
# cat events/sched/sched_switch/format
```

```
name: sched_switch
```

```
ID: 312
```

```
format:
```

```
    field:unsigned short common_type;    offset:0;    size:2;    signed:0;
    field:unsigned char common_flags;    offset:2;    size:1;    signed:0;
    field:unsigned char common_preempt_count; offset:3;    size:1;    signed:0;
    field:int common_pid;    offset:4;    size:4;    signed:1;
```

```
    field:char prev_comm[16];    offset:8;    size:16;    signed:1;
    field:pid_t prev_pid;    offset:24;    size:4;    signed:1;
    field:int prev_prio;    offset:28;    size:4;    signed:1;
    field:long prev_state;    offset:32;    size:8;    signed:1;
    field:char next_comm[16];    offset:40;    size:16;    signed:1;
    field:pid_t next_pid;    offset:56;    size:4;    signed:1;
    field:int next_prio;    offset:60;    size:4;    signed:1;
```

```
print fmt: "prev_comm=%s prev_pid=%d prev_prio=%d prev_state=%s%s ==> next_comm=%s next_pid=%d
next_prio=%d", REC->prev_comm, REC->prev_pid, REC->prev_prio, (REC->prev_state & (((0x0000 | 0x0001 |
0x0002 | 0x0004 | 0x0008 | 0x0010 | 0x0020 | 0x0040) + 1) << 1) - 1)) ? __print_flags(REC->prev_state
& (((0x0000 | 0x0001 | 0x0002 | 0x0004 | 0x0008 | 0x0010 | 0x0020 | 0x0040) + 1) << 1) - 1), "|",
{ 0x01, "S" }, { 0x02, "D" }, { 0x04, "T" }, { 0x08, "t" }, { 0x10, "X" }, { 0x20, "Z" }, { 0x40,
"P" }, { 0x80, "I" }) : "R", REC->prev_state & (((0x0000 | 0x0001 | 0x0002 | 0x0004 | 0x0008 | 0x0010
| 0x0020 | 0x0040) + 1) << 1) ? "+" : "", REC->next_comm, REC->next_pid, REC->next_prio
```

What's in an event?

```
# cat events/sched/sched_switch/format
```

```
name: sched_switch
```

```
ID: 312
```

```
format:
```

```
field:unsigned short common_type;   offset:0;   size:2;   signed:0;
field:unsigned char common_flags;   offset:2;   size:1;   signed:0;
field:unsigned char common_preempt_count; offset:3;   size:1;   signed:0;
field:int common_pid;   offset:4;   size:4;   signed:1;
```

```
field:char prev_comm[16];   offset:8;   size:16;   signed:1;
field:pid_t prev_pid;   offset:24;   size:4;   signed:1;
field:int prev_prio;   offset:28;   size:4;   signed:1;
field:long prev_state; offset:32;   size:8;   signed:1;
field:char next_comm[16];   offset:40;   size:16;   signed:1;
field:pid_t next_pid;   offset:56;   size:4;   signed:1;
field:int next_prio;   offset:60;   size:4;   signed:1;
```

```
print fmt: "prev_comm=%s prev_pid=%d prev_prio=%d prev_state=%s%s ==> next_comm=%s next_pid=%d
next_prio=%d", REC->prev_comm, REC->prev_pid, REC->prev_prio, (REC->prev_state & (((0x0000 | 0x0001 |
0x0002 | 0x0004 | 0x0008 | 0x0010 | 0x0020 | 0x0040) + 1) << 1) - 1) ? __print_flags(REC->prev_state
& (((0x0000 | 0x0001 | 0x0002 | 0x0004 | 0x0008 | 0x0010 | 0x0020 | 0x0040) + 1) << 1) - 1), "|",
{ 0x01, "S" }, { 0x02, "D" }, { 0x04, "T" }, { 0x08, "t" }, { 0x10, "X" }, { 0x20, "Z" }, { 0x40,
"P" }, { 0x80, "I" }) : "R", REC->prev_state & (((0x0000 | 0x0001 | 0x0002 | 0x0004 | 0x0008 | 0x0010
| 0x0020 | 0x0040) + 1) << 1) ? "+" : "", REC->next_comm, REC->next_pid, REC->next_prio
```

What's in an event?

```
# cat events/sched/sched_switch/format
```

```
name: sched_switch
```

```
ID: 312
```

```
format:
```

```
field:unsigned short common_type;   offset:0;   size:2;   signed:0;
field:unsigned char common_flags;   offset:2;   size:1;   signed:0;
field:unsigned char common_preempt_count; offset:3;   size:1;   signed:0;
field:int common_pid;   offset:4;   size:4;   signed:1;
```

```
field:char prev_comm[16];   offset:8;   size:16;   signed:1;
field:pid_t prev_pid;   offset:24;   size:4;   signed:1;
field:int prev_prio;   offset:28;   size:4;   signed:1;
field:long prev_state; offset:32;   size:8;   signed:1;
field:char next_comm[16];   offset:40;   size:16;   signed:1;
field:pid_t next_pid;   offset:56;   size:4;   signed:1;
field:int next_prio;   offset:60;   size:4;   signed:1;
```

```
print fmt: "prev_comm=%s prev_pid=%d prev_prio=%d prev_state=%s%s ==> next_comm=%s next_pid=%d
next_prio=%d", REC->prev_comm, REC->prev_pid, REC->prev_prio, (REC->prev_state & (((0x0000 | 0x0001 |
0x0002 | 0x0004 | 0x0008 | 0x0010 | 0x0020 | 0x0040) + 1) << 1) - 1)) ? __print_flags(REC->prev_state
& (((0x0000 | 0x0001 | 0x0002 | 0x0004 | 0x0008 | 0x0010 | 0x0020 | 0x0040) + 1) << 1) - 1), "|",
{ 0x01, "S" }, { 0x02, "D" }, { 0x04, "T" }, { 0x08, "t" }, { 0x10, "X" }, { 0x20, "Z" }, { 0x40,
"P" }, { 0x80, "I" }) : "R", REC->prev_state & (((0x0000 | 0x0001 | 0x0002 | 0x0004 | 0x0008 | 0x0010
| 0x0020 | 0x0040) + 1) << 1) ? "+" : "", REC->next_comm, REC->next_pid, REC->next_prio
```

What's in an event?

```
print fmt: "prev_comm=%s prev_pid=%d prev_prio=%d prev_state=%s%s ==> next_comm=%s next_pid=%d
next_prio=%d",
    REC->prev_comm, REC->prev_pid, REC->prev_prio,

    (REC->prev_state & (((0x0000 | 0x0001 | 0x0002 | 0x0004 | 0x0008 | 0x0010 | 0x0020 |
0x0040) + 1) << 1) - 1)) ?
    __print_flags(REC->prev_state & (((0x0000 | 0x0001 | 0x0002 | 0x0004 | 0x0008 | 0x0010 |
0x0020 | 0x0040) + 1) << 1) - 1), "|",
        { 0x01, "S" },
        { 0x02, "D" },
        { 0x04, "T" },
        { 0x08, "t" },
        { 0x10, "X" },
        { 0x20, "Z" },
        { 0x40, "P" },
        { 0x80, "I" }) :
    "R",

    REC->prev_state & (((0x0000 | 0x0001 | 0x0002 | 0x0004 | 0x0008 | 0x0010 | 0x0020 |
0x0040) + 1) << 1) ? "+" : "",
    REC->next_comm, REC->next_pid, REC->next_prio
```


What's in an event?

```
print fmt: "prev_comm=%s prev_pid=%d prev_prio=%d prev_state=%s%s ==> next_comm=%s next_pid=%d
next_prio=%d",
    REC->prev_comm, REC->prev_pid, REC->prev_prio,

    (REC->prev_state & (TASK_REPORT_MAX - 1)) ?
    __print_flags(REC->prev_state & (TASK_REPORT_MAX - 1), "|",
        { TASK_INTERRUPTIBLE, "S" },
        { TASK_UNINTERRUPTIBLE, "D" },
        { __TASK_STOPPED, "T" },
        { __TASK_TRACED, "t" },
        { EXIT_DEAD, "X" },
        { EXIT_ZOMBIE, "Z" },
        { TASK_PARKED, "P" },
        { TASK_DEAD, "I" }) :

    "R",

    REC->prev_state & TASK_REPORT_MAX ? "+" : "",
    REC->next_comm, REC->next_pid, REC->next_prio
```

Filtering Events

Too Much Info!

- This can be just as bad as not enough info
- Finding the needle in the haystack
- Signal to noise!

Filter out everything we do not want

- Just trace what we are interested in

The “filter” file

What's in an event?

```
# cat events/sched/sched_switch/format
```

```
name: sched_switch
```

```
ID: 312
```

```
format:
```

```
    field:unsigned short common_type;    offset:0;    size:2;    signed:0;
    field:unsigned char common_flags;    offset:2;    size:1;    signed:0;
    field:unsigned char common_preempt_count; offset:3;    size:1;    signed:0;
    field:int common_pid;    offset:4;    size:4;    signed:1;
```

```
    field:char prev_comm[16];    offset:8;    size:16;    signed:1;
    field:pid_t prev_pid;    offset:24;    size:4;    signed:1;
    field:int prev_prio;    offset:28;    size:4;    signed:1;
    field:long prev_state;    offset:32;    size:8;    signed:1;
    field:char next_comm[16];    offset:40;    size:16;    signed:1;
    field:pid_t next_pid;    offset:56;    size:4;    signed:1;
    field:int next_prio;    offset:60;    size:4;    signed:1;
```

Filtering Events

```
# echo 'prev_comm == "bash" && prev_state & 0x02' > events/sched/sched_switch/filter
# echo 1 > events/sched/sched_switch/enable
# echo > trace
# cat trace
```

```
# tracer: nop
```


##

		<pre> _-----> irqsoft /-----> need-resched /-----> hardirq/softirq /-----> preempt-depth /-----> delay </pre>										
TASK-PID	CPU#		TIMESTAMP	FUNCTION								
bash-27607	[000]	d...	1559578.742752:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/0	next_pid=0	next_prio=120
bash-27607	[001]	d...	1559579.026451:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/1	next_pid=0	next_prio=120
bash-27607	[000]	d...	1559579.111236:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/0	next_pid=0	next_prio=120
bash-27607	[000]	d...	1559579.171892:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/0	next_pid=0	next_prio=120
bash-27607	[002]	d...	1559579.304215:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=xfwm4	next_pid=17105	next_prio=120
bash-27607	[002]	d...	1559579.414473:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/2	next_pid=0	next_prio=120
bash-27607	[002]	d...	1559579.470549:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/2	next_pid=0	next_prio=120
bash-27607	[002]	d...	1559579.537746:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/2	next_pid=0	next_prio=120
bash-27607	[002]	d...	1559579.591502:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/2	next_pid=0	next_prio=120
bash-27607	[003]	d...	1559579.695324:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/3	next_pid=0	next_prio=120
bash-27607	[002]	d...	1559579.841453:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/2	next_pid=0	next_prio=120
bash-27607	[002]	d...	1559579.863734:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/2	next_pid=0	next_prio=120
bash-27607	[003]	d...	1559615.733229:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/3	next_pid=0	next_prio=120
bash-27607	[003]	d...	1559615.832504:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/3	next_pid=0	next_prio=120
bash-27607	[003]	d...	1559616.309440:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/3	next_pid=0	next_prio=120
bash-27607	[003]	d...	1559616.778076:	sched_switch:	prev_comm=bash	prev_pid=27607	prev_prio=120	prev_state=D	==>	next_comm=swapper/3	next_pid=0	next_prio=120

Event Triggers

Make an event do something special

- Turn off tracing
- Turn on tracing
- Take a “snapshot”
- Produce a stack dump
- Enable another event
- Disable another event

Filtered Events along with triggers

```
# echo 'prev_comm == "bash" && prev_state & 0x02' > events/sched/sched_switch/filter
# echo 'stacktrace if prev_comm == "bash" && prev_state & 0x02' > events/sched/sched_switch/trigger
# echo 1 > events/sched/sched_switch/enable
# echo > trace
# cat trace

# tracer: nop
#
#
#          _-----> irqs-off
#         /_-----> need-resched
#        |/_-----> hardirq/softirq
#       ||/_-----> preempt-depth
#      |||/_-----> delay
#     |||||/
#    TASK-PID  CPU#  |         |   TASK-PID  CPU#  |   FUNCTION
#    |||||-----|||   |||||-----|||   |||||-----|||
bash-27607 [002] d... 1559946.989729: sched_switch: prev_comm=bash prev_pid=27607 prev_prio=120 prev_state=D ==> next_comm=swapper/2 next_pid=0 next_prio=120
bash-27607 [002] d... 1559946.989746: <stack trace>
=> __schedule
=> schedule
=> schedule_hrtimeout_range_clock
=> poll_schedule_timeout.constprop.11
=> do_select
=> core_sys_select
=> do_pselect
=> __x64_sys_pselect6
=> do_syscall_64
=> entry_SYSCALL_64_after_hwframe
    bash-27607 [001] d... 1559947.145818: sched_switch: prev_comm=bash prev_pid=27607 prev_prio=120 prev_state=D ==> next_comm=swapper/1 next_pid=0 next_prio=120
    bash-27607 [001] d... 1559947.145833: <stack trace>
=> __schedule
=> schedule
=> schedule_hrtimeout_range_clock
=> poll_schedule_timeout.constprop.11
=> do_select
=> core_sys_select
=> do_pselect
```

Triggers are a little more difficult to remove

```
# echo '!stacktrace' > events/sched/sched_switch/trigger
```

fttrace - a tool for everyone with BusyBox

Everything we did so far used echo or cat

Makes using fttrace extremely simple on limited systems

But can be very tedious

- Lots of things to remember (what file does what)
- Not very intuitive
- Can be painstaking
- Hard to do batch processing
- Hard to record specific functions

Introducing trace-cmd

You did make that clone didn't you?

Is an executable that interacts with the ftrace interface

No need to worry about the tracefs system

- It does it for you

Can also save the data to a file

- Uses per_cpu/cpuX/trace_pipe_raw
- Reads the binary data directly from the ring buffer
- Uses splice(2) to write directly to a file or network (zero copy)

Uses the format files to know how to read the binary data

Introducing trace-cmd

Forget everything you learned so far

- (no don't really, but let's start over)

From now on, be in a directory that you can write to

- Still be root user, or at least start all commands with “sudo”

trace-cmd will do the work for you

If you did “make install_doc”

- man trace-cmd
- man trace-cmd record
- man trace-cmd report
- etc

trace-cmd start and show

```
# trace-cmd start -e sched_switch -f 'prev_comm == "bash" && prev_state & 0x02' \
-R 'stracktrace if prev_comm == "bash" && prev_state & 0x02'
# trace-cmd show

# tracer: nop
#
#          _-----> irqs-off
#         / _-----> need-resched
#        | / _-----> hardirq/softirq
#       || / _-----> preempt-depth
#      ||| / _-----> delay
#     |||| /
# TASK-PID  CPU#  ||||  TIMESTAMP  FUNCTION
#   |||       |||       |||          |
bash-7855  [000] d... 1564606.264976: sched_switch: prev_comm=bash prev_pid=7855 prev_prio=120 prev_state=D ==> next_comm=swapper/0 next_pid=0 next_prio=120
bash-7855  [000] d... 1564606.264994: <stack trace>
=> __schedule
=> schedule
=> schedule_hrtimeout_range_clock
=> poll_schedule_timeout.constprop.11
=> do_select
=> core_sys_select
=> do_pselect
=> __x64_sys_pselect6
=> do_syscall_64
=> entry_SYSCALL_64_after_hwframe
bash-7855  [001] d... 1564606.664732: sched_switch: prev_comm=bash prev_pid=7855 prev_prio=120 prev_state=D ==> next_comm=SendControllerT next_pid=12716
next_prio=120
bash-7855  [001] d... 1564606.664749: <stack trace>
=> __schedule
=> schedule
=> schedule_hrtimeout_range_clock
=> poll_schedule_timeout.constprop.11
=> do_select
=> core_sys_select
=> do_pselect
```

trace-cmd stat and reset

```
# trace-cmd stat

Events:
  Individual events:
    sched
    sched_switch

Filters:
  sched:sched_switch "(prev_comm == "bash" && prev_state & 0x02)"

Triggers:
  sched:sched_switch "stacktrace:unlimited if prev_comm == "bash" && prev_state & 0x02"

Buffer size in kilobytes (per cpu):
  1408

Buffer total size in kilobytes:
  5632

Tracing is enabled

# trace-cmd reset
# trace-cmd stat

Events:
  All disabled

Buffer size in kilobytes (per cpu):
  1408

Buffer total size in kilobytes:
  5632

Tracing is disabled
```

trace-cmd start -p nop

```
# trace-cmd reset
# trace-cmd stat

Events:
  All disabled

Buffer size in kilobytes (per cpu):
  1408

Buffer total size in kilobytes:
  5632

Tracing is disabled

# trace-cmd start -p nop
# trace-cmd stat

Events:
  All disabled

Buffer size in kilobytes (per cpu):
  1408

Buffer total size in kilobytes:
  5632

Tracing is enabled
```

trace-cmd record and start

“start” enables tracing but does no recording

- Use “show” to see the in-kernel ring buffer output “trace” file.

“record” records the trace data into a file (default: trace.dat)

- Use “report” to read the file
- The “record” will zero copy from the per_cpu files

“start” has most the same options as “record”

- To enable tracing
- “-e” for events
- “-p” for tracers (historically they were once called “plugins”)

Now let's start seeing what your computer is doing

Isn't that the title of this tutorial?

We spent enough time on details, let's start doing something fun

Let's write a bash script

And see exactly what goes on in the kernel!

Make a really simple program

Using your favorite editor, create this file

```
#!/bin/bash  
  
echo "Hello world!"
```


Introduction (or not) to strace

strace - Traces the system calls a program makes

- See how it communicates with the kernel
- It understands the calls that are made
- Shows file names and parameters of the system calls

We will start with this before jumping into the kernel

strace our Hello World!

```
# strace ./hello &> out
# cat out

execve("./hello", [ "./hello" ], 0x7ffebab3e880 /* 13 vars */) = 0
brk(NULL)                                = 0x559b65f84000
access("/etc/ld.so.nohwcap", F_OK)        = -1 ENOENT (No such file or directory)
access("/etc/ld.so.preload", R_OK)        = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=156426, ...}) = 0
mmap(NULL, 156426, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f446db12000
close(3)                                  = 0
access("/etc/ld.so.nohwcap", F_OK)        = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libtinfo.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\0\351\0\0\0\0\0"..., 832) = 832
fstat(3, {st_mode=S_IFREG|0644, st_size=183528, ...}) = 0
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f446db10000
[ more mmap calls ]
close(3)                                  = 0
access("/etc/ld.so.nohwcap", F_OK)        = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libdl.so.2", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\0\000\21\0\0\0\0\0"..., 832) = 832
fstat(3, {st_mode=S_IFREG|0644, st_size=14592, ...}) = 0
[ more mmap calls ]
close(3)                                  = 0
access("/etc/ld.so.nohwcap", F_OK)        = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\0\260A\2\0\0\0\0\0"..., 832) = 832
fstat(3, {st_mode=S_IFREG|0755, st_size=1824496, ...}) = 0
[ more mmap calls ]
close(3)                                  = 0
```

strace our Hello World!

```
# strace ./hello &> out
# cat out

[...]
```

```
openat(AT_FDCWD, "./hello", O_RDONLY) = 3
stat("./hello", {st_mode=S_IFREG|0755, st_size=33, ...}) = 0
ioctl(3, TCGETS, 0x7ffcfa6a6950) = -1 ENOTTY (Inappropriate ioctl for device)
lseek(3, 0, SEEK_CUR) = 0
read(3, "#!/bin/bash\n\nnecho 'hello world!'"..., 80) = 33
lseek(3, 0, SEEK_SET) = 0
prlimit64(0, RLIMIT_NOFILE, NULL, {rlim_cur=1024, rlim_max=1024*1024}) = 0
fcntl(255, F_GETFD) = -1 EBADF (Bad file descriptor)
dup2(3, 255) = 255
close(3) = 0
fcntl(255, F_SETFD, FD_CLOEXEC) = 0
fcntl(255, F_GETFL) = 0x8000 (flags O_RDONLY|O_LARGEFILE)
fstat(255, {st_mode=S_IFREG|0755, st_size=33, ...}) = 0
lseek(255, 0, SEEK_CUR) = 0
read(255, "#!/bin/bash\n\nnecho 'hello world!'"..., 33) = 33
fstat(1, {st_mode=S_IFCHR|0600, st_rdev=makedev(136, 3), ...}) = 0
write(1, "hello world!\n", 13) = 13
read(255, "", 33) = 0
rt_sigprocmask(SIG_BLOCK, [CHLD], [], 8) = 0
rt_sigprocmask(SIG_SETMASK, [], NULL, 8) = 0
exit_group(0) = ?
+++ exited with 0 +++
```

strace our Hello World!

```
# strace ./hello &> out
# cat out

[...]
```

```
openat(AT_FDCWD, "./hello", O_RDONLY) = 3
stat("./hello", {st_mode=S_IFREG|0755, st_size=33, ...}) = 0
ioctl(3, TCGETS, 0x7ffcfa6a6950) = -1 ENOTTY (Inappropriate ioctl for device)
lseek(3, 0, SEEK_CUR) = 0
read(3, "#!/bin/bash\n\nnecho 'hello world!'"..., 80) = 33
lseek(3, 0, SEEK_SET) = 0
prlimit64(0, RLIMIT_NOFILE, NULL, {rlim_cur=1024, rlim_max=1024*1024}) = 0
fcntl(255, F_GETFD) = -1 EBADF (Bad file descriptor)
dup2(3, 255) = 255
close(3) = 0
fcntl(255, F_SETFD, FD_CLOEXEC) = 0
fcntl(255, F_GETFL) = 0x8000 (flags O_RDONLY|O_LARGEFILE)
fstat(255, {st_mode=S_IFREG|0755, st_size=33, ...}) = 0
lseek(255, 0, SEEK_CUR) = 0
read(255, "#!/bin/bash\n\nnecho 'hello world!'"..., 33) = 33
fstat(1, {st_mode=S_IFCHR|0600, st_rdev=makedev(136, 3), ...}) = 0
write(1, "hello world!\n", 13hello world!
) = 13
read(255, "", 33) = 0
rt_sigprocmask(SIG_BLOCK, [CHLD], [], 8) = 0
rt_sigprocmask(SIG_SETMASK, [], NULL, 8) = 0
exit_group(0) = ?
+++ exited with 0 +++
```

How do bash scripts work?

```
# strace ./hello &> out
# cat out

[..]
execve("./hello", ["./hello"], 0x7ffebab3e880 /* 13 vars */) = 0

[..]
openat(AT_FDCWD, "./hello", O_RDONLY) = 3

[..]
read(255, "#!/bin/bash\n\nnecho 'hello world!'"..., 33) = 33

[..]
write(1, "hello world!\n", 13) = 13
```

Let's have a look at execve

In the kernel, system calls start with arch specific headers

__x64_sys_execve “__x64_sys_” is added to the system call name

```
# trace-cmd list -f execve
```

```
audit_log_execve_info  
__do_execve_file.isra.35  
__ia32_compat_sys_execve  
__ia32_compat_sys_execveat  
__ia32_sys_execve  
__ia32_sys_execveat  
__x32_compat_sys_execve  
__x64_sys_execve  
__x64_sys_execveat  
__x32_compat_sys_execveat  
do_execve_file  
do_execve  
do_execveat
```

trace-cmd list -f execve?

For more information, do: “man trace-cmd list”

trace-cmd list -f execve?

For more information, do: “man trace-cmd list”

Moving on, we don't have time for details!

trace-cmd list -f execve?

For more information, do: “man trace-cmd list”

Moving on, we don't have time for details!

All you need to know, is that we found our execve system call function

Graph our Hello World!

After all that, I cheat with “*sys_exec*”!

```
# trace-cmd record -p function_graph -g '*sys_exec*' ./hello

  plugin 'function_graph'
hello world!
CPU0 data recorded at offset=0x652000
    0 bytes in size
CPU1 data recorded at offset=0x652000
  851968 bytes in size
CPU2 data recorded at offset=0x722000
    0 bytes in size
CPU3 data recorded at offset=0x722000
  12288 bytes in size
```

Graph our Hello World!

trace-cmd report

CPU 0 is empty
CPU 2 is empty
cpus=4

```
hello-16873 [003] 1653533.226805: funcgraph_entry:
hello-16873 [003] 1653533.226808: funcgraph_entry:
hello-16873 [003] 1653533.226808: funcgraph_entry:
hello-16873 [003] 1653533.226809: funcgraph_entry:
hello-16873 [003] 1653533.226809: funcgraph_entry:
hello-16873 [003] 1653533.226809: funcgraph_entry: 0.054 us
hello-16873 [003] 1653533.226809: funcgraph_exit: 0.431 us
hello-16873 [003] 1653533.226809: funcgraph_entry: 0.032 us
hello-16873 [003] 1653533.226810: funcgraph_entry: 0.038 us
hello-16873 [003] 1653533.226810: funcgraph_entry: 0.045 us
hello-16873 [003] 1653533.226810: funcgraph_exit: 1.611 us
hello-16873 [003] 1653533.226810: funcgraph_entry:
hello-16873 [003] 1653533.226811: funcgraph_entry: 0.062 us
hello-16873 [003] 1653533.226811: funcgraph_entry: 0.043 us
hello-16873 [003] 1653533.226811: funcgraph_entry: 0.042 us
hello-16873 [003] 1653533.226812: funcgraph_exit: 1.023 us
hello-16873 [003] 1653533.226812: funcgraph_exit: 3.253 us
hello-16873 [003] 1653533.226812: funcgraph_exit: 3.583 us
hello-16873 [003] 1653533.226812: funcgraph_entry:
hello-16873 [003] 1653533.226812: funcgraph_entry:
hello-16873 [003] 1653533.226812: funcgraph_entry: 0.048 us
hello-16873 [003] 1653533.226813: funcgraph_exit: 0.354 us
hello-16873 [003] 1653533.226813: funcgraph_entry:
hello-16873 [003] 1653533.226813: funcgraph_entry:
hello-16873 [003] 1653533.226813: funcgraph_entry: 0.032 us
hello-16873 [003] 1653533.226813: funcgraph_exit: 0.310 us
hello-16873 [003] 1653533.226814: funcgraph_entry: 0.030 us
hello-16873 [003] 1653533.226814: funcgraph_entry: 0.106 us
hello-16873 [003] 1653533.226814: funcgraph_entry: 0.119 us
hello-16873 [003] 1653533.226815: funcgraph_exit: 1.845 us
```

```
__x64_sys_execve() {
  getname() {
    getname_flags() {
      kmem_cache_alloc() {
        _cond_resched() {
          rcu_all_qs();
        }
      }
      should_failslab();
      prefetch_freepointer();
      memcg_kmem_put_cache();
    }
    _check_object_size() {
      __virt_addr_valid();
      __check_heap_object();
      check_stack_object();
    }
  }
}
__do_execve_file.isra.35() {
  unshare_files() {
    unshare_fd();
  }
  kmem_cache_alloc_trace() {
    _cond_resched() {
      rcu_all_qs();
    }
  }
  should_failslab();
  prefetch_freepointer();
  memcg_kmem_put_cache();
}
```

There's a lot of info here!

```
# trace-cmd report | wc -l  
16870
```

There's a lot of info here!

```
# trace-cmd report | wc -l  
16870
```

Interrupts do happen

```
# trace-cmd report -I | wc -l  
10196
```

There's a lot of info here!

```
# trace-cmd report | wc -l  
16870
```

Interrupts do happen

```
# trace-cmd report -I | wc -l  
10196
```

As well as “soft” interrupts

```
# trace-cmd report -IS | wc -l  
10092
```

Only trace a few functions down

```
# trace-cmd record -p function_graph -g '*sys_exec*' --max-graph-depth 5 ./hello

  plugin 'function_graph'
hello world!
CPU0 data recorded at offset=0x652000
    4096 bytes in size
CPU1 data recorded at offset=0x653000
    0 bytes in size
CPU2 data recorded at offset=0x653000
    0 bytes in size
CPU3 data recorded at offset=0x653000
   16384 bytes in size
```

More likely prevents interrupts from being recorded

```
# trace-cmd report | wc -l  
347  
# trace-cmd report -I | wc -l  
347  
# trace-cmd report -IS | wc -l  
347
```


Remove “_cond_resched” calls (too many)

```
# trace-cmd record -p function_graph -g '*sys_exec*' --max-graph-depth 5 \  
-n _cond_resched ./hello
```

```
plugin 'function_graph'  
hello world!  
CPU0 data recorded at offset=0x652000  
4096 bytes in size  
CPU1 data recorded at offset=0x653000  
0 bytes in size  
CPU2 data recorded at offset=0x653000  
0 bytes in size  
CPU3 data recorded at offset=0x653000  
12288 bytes in size
```

Much better

```
# trace-cmd report | wc -l  
244  
# trace-cmd report -I | wc -l  
242  
# trace-cmd report -IS | wc -l  
242
```

And let's just do 4 deep

```
# trace-cmd record -p function_graph -g '*sys_exec*' --max-graph-depth 4 \  
-n _cond_resched ./hello
```

```
plugin 'function_graph'  
hello world!  
CPU0 data recorded at offset=0x652000  
0 bytes in size  
CPU1 data recorded at offset=0x652000  
4096 bytes in size  
CPU2 data recorded at offset=0x653000  
0 bytes in size  
CPU3 data recorded at offset=0x653000  
8192 bytes in size
```

Easier to understand

```
# trace-cmd report -IS
```

```
CPU 0 is empty  
CPU 2 is empty  
cpus=4
```

hello-18623	[001]	1654799.238555:	funcgraph_entry:		__x64_sys_execve() {
hello-18623	[001]	1654799.238561:	funcgraph_entry:		getname() {
hello-18623	[001]	1654799.238561:	funcgraph_entry:		getname_flags() {
hello-18623	[001]	1654799.238562:	funcgraph_entry:	0.373 us	kmem_cache_alloc();
hello-18623	[001]	1654799.238563:	funcgraph_entry:	0.232 us	__check_object_size();
hello-18623	[001]	1654799.238563:	funcgraph_exit:	1.613 us	}
hello-18623	[001]	1654799.238563:	funcgraph_exit:	2.040 us	}
hello-18623	[001]	1654799.238564:	funcgraph_entry:		__do_execve_file.isra.35() {
hello-18623	[001]	1654799.238564:	funcgraph_entry:		smp_irq_work_interrupt() {
hello-18623	[001]	1654799.238565:	funcgraph_entry:		irq_enter() {
hello-18623	[001]	1654799.238565:	funcgraph_entry:	0.094 us	rcu_irq_enter();
hello-18623	[001]	1654799.238572:	funcgraph_entry:	0.053 us	idle_cpu();
hello-18623	[001]	1654799.238573:	funcgraph_entry:	0.109 us	rcu_irq_exit();
hello-18623	[001]	1654799.238573:	funcgraph_exit:	1.058 us	}
hello-18623	[001]	1654799.238574:	funcgraph_exit:	8.978 us	}
hello-18623	[001]	1654799.238574:	funcgraph_entry:		unshare_files() {
hello-18623	[001]	1654799.238575:	funcgraph_entry:	0.079 us	unshare_fd();
hello-18623	[001]	1654799.238575:	funcgraph_exit:	0.625 us	}
hello-18623	[001]	1654799.238575:	funcgraph_entry:		kmem_cache_alloc_trace() {
hello-18623	[001]	1654799.238576:	funcgraph_entry:	0.057 us	should_failslab();
hello-18623	[001]	1654799.238577:	funcgraph_entry:	0.065 us	prefetch_freepointer();
hello-18623	[001]	1654799.238577:	funcgraph_entry:	0.064 us	memcg_kmem_put_cache();
hello-18623	[001]	1654799.238578:	funcgraph_exit:	2.141 us	}
hello-18623	[001]	1654799.238578:	funcgraph_entry:		prepare_bprm_creds() {
hello-18623	[001]	1654799.238578:	funcgraph_entry:	0.247 us	mutex_lock_interruptible();
hello-18623	[001]	1654799.238579:	funcgraph_entry:	1.410 us	prepare_exec_creds();
hello-18623	[001]	1654799.238580:	funcgraph_exit:	2.375 us	}
hello-18623	[001]	1654799.238581:	funcgraph_entry:	0.049 us	_raw_spin_lock();

Easier to understand

```
# trace-cmd report -IS
```

```
CPU 0 is empty  
CPU 2 is empty  
cpus=4
```

hello-18623	[001]	1654799.238555:	funcgraph_entry:		__x64_sys_execve() {
hello-18623	[001]	1654799.238561:	funcgraph_entry:		getname() {
hello-18623	[001]	1654799.238561:	funcgraph_entry:		getname_flags() {
hello-18623	[001]	1654799.238562:	funcgraph_entry:	0.373 us	kmem_cache_alloc();
hello-18623	[001]	1654799.238563:	funcgraph_entry:	0.232 us	__check_object_size();
hello-18623	[001]	1654799.238563:	funcgraph_exit:	1.613 us	}
hello-18623	[001]	1654799.238563:	funcgraph_exit:	2.040 us	}
hello-18623	[001]	1654799.238564:	funcgraph_entry:		__do_execve_file.isra.35() {
hello-18623	[001]	1654799.238564:	funcgraph_entry:		smp_irq_work_interrupt() {
hello-18623	[001]	1654799.238565:	funcgraph_entry:		irq_enter() {
hello-18623	[001]	1654799.238565:	funcgraph_entry:	0.094 us	rcu_irq_enter();
hello-18623	[001]	1654799.238572:	funcgraph_entry:	0.053 us	idle_cpu();
hello-18623	[001]	1654799.238573:	funcgraph_entry:	0.109 us	rcu_irq_exit();
hello-18623	[001]	1654799.238573:	funcgraph_exit:	1.058 us	}
hello-18623	[001]	1654799.238574:	funcgraph_exit:	8.978 us	}
hello-18623	[001]	1654799.238574:	funcgraph_entry:		unshare_files() {
hello-18623	[001]	1654799.238575:	funcgraph_entry:	0.079 us	unshare_fd();
hello-18623	[001]	1654799.238575:	funcgraph_exit:	0.625 us	}
hello-18623	[001]	1654799.238575:	funcgraph_entry:		kmem_cache_alloc_trace() {
hello-18623	[001]	1654799.238576:	funcgraph_entry:	0.057 us	should_failslab();
hello-18623	[001]	1654799.238577:	funcgraph_entry:	0.065 us	prefetch_freepointer();
hello-18623	[001]	1654799.238577:	funcgraph_entry:	0.064 us	memcg_kmem_put_cache();
hello-18623	[001]	1654799.238578:	funcgraph_exit:	2.141 us	}
hello-18623	[001]	1654799.238578:	funcgraph_entry:		prepare_bprm_creds() {
hello-18623	[001]	1654799.238578:	funcgraph_entry:	0.247 us	mutex_lock_interruptible();
hello-18623	[001]	1654799.238579:	funcgraph_entry:	1.410 us	prepare_exec_creds();
hello-18623	[001]	1654799.238580:	funcgraph_exit:	2.375 us	}
hello-18623	[001]	1654799.238581:	funcgraph_entry:	0.049 us	_raw_spin_lock();

Easier to understand

```
hello-18623 1.... 1654799.238581: funcgraph_entry:
hello-18623 1.... 1654799.238581: funcgraph_entry:      3.846 us
hello-18623 1.... 1654799.238585: funcgraph_entry:      0.035 us
hello-18623 1.... 1654799.238586: funcgraph_entry:      0.041 us
hello-18623 1.... 1654799.238586: funcgraph_exit:       4.887 us
hello-18623 1.... 1654799.238586: funcgraph_entry:
hello-18623 1.... 1654799.238586: funcgraph_entry:      0.036 us
hello-18623 1d... 1654799.238587: funcgraph_entry:      1.016 us
hello-18623 1d... 1654799.238588: funcgraph_entry:      0.044 us
hello-18623 1.... 1654799.238588: funcgraph_entry:      + 11.455 us
hello-18623 3d... 1654799.238602: funcgraph_entry:      1.061 us
hello-18623 3.... 1654799.238603: funcgraph_exit:       + 17.049 us
hello-18623 3.... 1654799.238604: funcgraph_entry:
hello-18623 3.... 1654799.238604: funcgraph_entry:      0.809 us
hello-18623 3.... 1654799.238605: funcgraph_entry:      2.296 us
hello-18623 3.... 1654799.238608: funcgraph_exit:      3.806 us
hello-18623 3.... 1654799.238608: funcgraph_entry:      0.041 us
hello-18623 3.... 1654799.238608: funcgraph_entry:
hello-18623 3.... 1654799.238608: funcgraph_entry:      0.706 us
hello-18623 3.... 1654799.238609: funcgraph_exit:      1.001 us
hello-18623 3.... 1654799.238609: funcgraph_entry:      0.222 us
hello-18623 3.... 1654799.238610: funcgraph_entry:      0.045 us
hello-18623 3.... 1654799.238610: funcgraph_entry:
hello-18623 3.... 1654799.238610: funcgraph_entry:      0.310 us
hello-18623 3.... 1654799.238611: funcgraph_entry:      0.347 us
hello-18623 3.... 1654799.238612: funcgraph_exit:      1.295 us
hello-18623 3.... 1654799.238612: funcgraph_entry:      0.037 us
hello-18623 3.... 1654799.238612: funcgraph_entry:      0.382 us
hello-18623 3.... 1654799.238613: funcgraph_entry:      1.817 us
hello-18623 3.... 1654799.238615: funcgraph_entry:
hello-18623 3.... 1654799.238615: funcgraph_entry:      0.216 us
hello-18623 3.... 1654799.238615: funcgraph_entry:      9.257 us
hello-18623 3.... 1654799.238625: funcgraph_entry:      2.590 us
hello-18623 3.... 1654799.238628: funcgraph_exit:      + 12.984 us
```

```
do_open_execat() {
    do_filp_open();
    __fsnotify_parent();
    fsnotify();
}
sched_exec() {
    _raw_spin_lock_irqsave();
    select_task_rq_fair();
    _raw_spin_unlock_irqrestore();
    stop_one_cpu();
    smp_irq_work_interrupt();
}
mm_alloc() {
    kmem_cache_alloc();
    mm_init();
}
_raw_spin_lock();
vm_area_alloc() {
    kmem_cache_alloc();
}
down_write_killable();
vm_get_page_prot();
insert_vm_struct() {
    security_vm_enough_memory_mm();
    vma_link();
}
up_write();
count.isra.24.constprop.38();
count.isra.24.constprop.38();
prepare_binprm() {
    mnt_may_suid();
    security_bprm_set_creds();
    kernel_read();
}
```

Easier to understand

hello-18623	3....	1654799.238628:	funcgraph_entry:			copy_strings_kernel() {
hello-18623	3....	1654799.238628:	funcgraph_entry:	+ 10.840 us		copy_strings.isra.26();
hello-18623	3....	1654799.238639:	funcgraph_exit:	+ 11.232 us		}
hello-18623	3....	1654799.238640:	funcgraph_entry:			copy_strings.isra.26() {
hello-18623	3....	1654799.238640:	funcgraph_entry:	0.549 us		get_user_pages_remote();
hello-18623	3....	1654799.238641:	funcgraph_entry:	0.090 us		__check_object_size();
[...]						
hello-18623	3....	1654799.238648:	funcgraph_entry:	0.085 us		__check_object_size();
hello-18623	3....	1654799.238649:	funcgraph_exit:	8.993 us		}
hello-18623	3....	1654799.238649:	funcgraph_entry:			copy_strings.isra.26() {
hello-18623	3....	1654799.238649:	funcgraph_entry:	0.449 us		get_user_pages_remote();
hello-18623	3....	1654799.238650:	funcgraph_entry:	0.089 us		__check_object_size();
hello-18623	3....	1654799.238650:	funcgraph_exit:	1.444 us		}
hello-18623	3....	1654799.238650:	funcgraph_entry:			would_dump() {
hello-18623	3....	1654799.238651:	funcgraph_entry:	0.119 us		inode_permission();
hello-18623	3....	1654799.238651:	funcgraph_exit:	0.415 us		}
hello-18623	3....	1654799.238651:	funcgraph_entry:	0.114 us		task_active_pid_ns();
hello-18623	3....	1654799.238652:	funcgraph_entry:	0.079 us		__task_pid_nr_ns();
hello-18623	3....	1654799.238652:	funcgraph_entry:			search_binary_handler() {
hello-18623	3....	1654799.238652:	funcgraph_entry:	0.048 us		security_bprm_check();
hello-18623	3....	1654799.238652:	funcgraph_entry:	0.057 us		_raw_read_lock();
hello-18623	3....	1654799.238653:	funcgraph_entry:	0.036 us		try_module_get();
hello-18623	3....	1654799.238653:	funcgraph_entry:	! 214.953 us		load_script();
hello-18623	3....	1654799.238869:	funcgraph_entry:	0.052 us		_raw_read_lock();
hello-18623	3....	1654799.238869:	funcgraph_entry:	0.034 us		module_put();
hello-18623	3....	1654799.238869:	funcgraph_exit:	! 217.327 us		}
hello-18623	3....	1654799.238870:	funcgraph_entry:	0.066 us		proc_exec_connector();
hello-18623	3....	1654799.238870:	funcgraph_entry:			acct_update_integrals() {
hello-18623	3d...	1654799.238870:	funcgraph_entry:	0.042 us		__acct_update_integrals();
hello-18623	3....	1654799.238870:	funcgraph_exit:	0.336 us		}
hello-18623	3....	1654799.238871:	funcgraph_entry:			task_numa_free() {
hello-18623	3....	1654799.238871:	funcgraph_entry:	0.040 us		kfree();
hello-18623	3d...	1654799.238872:	funcgraph_entry:	3.233 us		smp_irq_work_interrupt();
hello-18623	3....	1654799.238876:	funcgraph_exit:	5.134 us		}

Easier to understand

hello-18623	3....	1654799.238628:	funcgraph_entry:			copy_strings_kernel() {
hello-18623	3....	1654799.238628:	funcgraph_entry:	+ 10.840 us		copy_strings.isra.26();
hello-18623	3....	1654799.238639:	funcgraph_exit:	+ 11.232 us		}
hello-18623	3....	1654799.238640:	funcgraph_entry:			copy_strings.isra.26() {
hello-18623	3....	1654799.238640:	funcgraph_entry:	0.549 us		get_user_pages_remote();
hello-18623	3....	1654799.238641:	funcgraph_entry:	0.090 us		__check_object_size();
[..]						
hello-18623	3....	1654799.238648:	funcgraph_entry:	0.085 us		__check_object_size();
hello-18623	3....	1654799.238649:	funcgraph_exit:	8.993 us		}
hello-18623	3....	1654799.238649:	funcgraph_entry:			copy_strings.isra.26() {
hello-18623	3....	1654799.238649:	funcgraph_entry:	0.449 us		get_user_pages_remote();
hello-18623	3....	1654799.238650:	funcgraph_entry:	0.089 us		__check_object_size();
hello-18623	3....	1654799.238650:	funcgraph_exit:	1.444 us		}
hello-18623	3....	1654799.238650:	funcgraph_entry:			would_dump() {
hello-18623	3....	1654799.238651:	funcgraph_entry:	0.119 us		inode_permission();
hello-18623	3....	1654799.238651:	funcgraph_exit:	0.415 us		}
hello-18623	3....	1654799.238651:	funcgraph_entry:	0.114 us		task_active_pid_ns();
hello-18623	3....	1654799.238652:	funcgraph_entry:	0.079 us		__task_pid_nr_ns();
hello-18623	3....	1654799.238652:	funcgraph_entry:			search_binary_handler() {
hello-18623	3....	1654799.238652:	funcgraph_entry:	0.048 us		security_bprm_check();
hello-18623	3....	1654799.238652:	funcgraph_entry:	0.057 us		_raw_read_lock();
hello-18623	3....	1654799.238653:	funcgraph_entry:	0.036 us		try_module_get();
hello-18623	3....	1654799.238653:	funcgraph_entry:	! 214.953 us		load_script();
hello-18623	3....	1654799.238869:	funcgraph_entry:	0.052 us		_raw_read_lock();
hello-18623	3....	1654799.238869:	funcgraph_entry:	0.034 us		module_put();
hello-18623	3....	1654799.238869:	funcgraph_exit:	! 217.327 us		}
hello-18623	3....	1654799.238870:	funcgraph_entry:	0.066 us		proc_exec_connector();
hello-18623	3....	1654799.238870:	funcgraph_entry:			acct_update_integrals() {
hello-18623	3d...	1654799.238870:	funcgraph_entry:	0.042 us		__acct_update_integrals();
hello-18623	3....	1654799.238870:	funcgraph_exit:	0.336 us		}
hello-18623	3....	1654799.238871:	funcgraph_entry:			task_numa_free() {
hello-18623	3....	1654799.238871:	funcgraph_entry:	0.040 us		kfree();
hello-18623	3d...	1654799.238872:	funcgraph_entry:	3.233 us		smp_irq_work_interrupt();
hello-18623	3....	1654799.238876:	funcgraph_exit:	5.134 us		}

Easier to understand

hello-18623	3....	1654799.238628:	funcgraph_entry:			copy_strings_kernel() {
hello-18623	3....	1654799.238628:	funcgraph_entry:	+ 10.840 us		copy_strings.isra.26();
hello-18623	3....	1654799.238639:	funcgraph_exit:	+ 11.232 us		}
hello-18623	3....	1654799.238640:	funcgraph_entry:			copy_strings.isra.26() {
hello-18623	3....	1654799.238640:	funcgraph_entry:	0.549 us		get_user_pages_remote();
hello-18623	3....	1654799.238641:	funcgraph_entry:	0.090 us		__check_object_size();
[...]						
hello-18623	3....	1654799.238648:	funcgraph_entry:	0.085 us		__check_object_size();
hello-18623	3....	1654799.238649:	funcgraph_exit:	8.993 us		}
hello-18623	3....	1654799.238649:	funcgraph_entry:			copy_strings.isra.26() {
hello-18623	3....	1654799.238649:	funcgraph_entry:	0.449 us		get_user_pages_remote();
hello-18623	3....	1654799.238650:	funcgraph_entry:	0.089 us		__check_object_size();
hello-18623	3....	1654799.238650:	funcgraph_exit:	1.444 us		}
hello-18623	3....	1654799.238650:	funcgraph_entry:			would_dump() {
hello-18623	3....	1654799.238651:	funcgraph_entry:	0.119 us		inode_permission();
hello-18623	3....	1654799.238651:	funcgraph_exit:	0.415 us		}
hello-18623	3....	1654799.238651:	funcgraph_entry:	0.114 us		task_active_pid_ns();
hello-18623	3....	1654799.238652:	funcgraph_entry:	0.079 us		__task_pid_nr_ns();
hello-18623	3....	1654799.238652:	funcgraph_entry:			search_binary_handler() {
hello-18623	3....	1654799.238652:	funcgraph_entry:	0.048 us		security_bprm_check();
hello-18623	3....	1654799.238652:	funcgraph_entry:	0.057 us		_raw_read_lock();
hello-18623	3....	1654799.238653:	funcgraph_entry:	0.036 us		try_module_get();
hello-18623	3....	1654799.238653:	funcgraph_entry:	! 214.953 us		load_script();
hello-18623	3....	1654799.238869:	funcgraph_entry:	0.052 us		_raw_read_lock();
hello-18623	3....	1654799.238869:	funcgraph_entry:	0.034 us		module_put();
hello-18623	3....	1654799.238869:	funcgraph_exit:	! 217.327 us		}
hello-18623	3....	1654799.238870:	funcgraph_entry:	0.066 us		proc_exec_connector();
hello-18623	3....	1654799.238870:	funcgraph_entry:			acct_update_integrals() {
hello-18623	3d...	1654799.238870:	funcgraph_entry:	0.042 us		__acct_update_integrals();
hello-18623	3....	1654799.238870:	funcgraph_exit:	0.336 us		}
hello-18623	3....	1654799.238871:	funcgraph_entry:			task_numa_free() {
hello-18623	3....	1654799.238871:	funcgraph_entry:	0.040 us		kfree();
hello-18623	3d...	1654799.238872:	funcgraph_entry:	3.233 us		smp_irq_work_interrupt();
hello-18623	3....	1654799.238876:	funcgraph_exit:	5.134 us		}

What's this load_script doing?

```
# trace-cmd record -p function_graph -g 'load_script' --max-graph-depth 3 \  
-n _cond_resched ./hello
```

```
plugin 'function_graph'  
hello world!  
CPU0 data recorded at offset=0x652000  
    0 bytes in size  
CPU1 data recorded at offset=0x652000  
    0 bytes in size  
CPU2 data recorded at offset=0x652000  
    0 bytes in size  
CPU3 data recorded at offset=0x652000  
    4096 bytes in size
```

Easier to understand

```
# trace-cmd report -IS
```

```
CPU 0 is empty
CPU 1 is empty
CPU 2 is empty
cpus=4
```

hello-20414	3....	1656091.711043:	funcgraph_entry:		load_script() {
hello-20414	3....	1656091.711046:	funcgraph_entry:		fput() {
hello-20414	3....	1656091.711046:	funcgraph_entry:	0.143 us	task_work_add();
hello-20414	3....	1656091.711047:	funcgraph_exit:	0.630 us	}
hello-20414	3....	1656091.711047:	funcgraph_entry:		remove_arg_zero() {
hello-20414	3....	1656091.711047:	funcgraph_entry:	0.625 us	get_user_pages_remote();
hello-20414	3....	1656091.711048:	funcgraph_exit:	0.971 us	}
hello-20414	3....	1656091.711048:	funcgraph_entry:		copy_strings_kernel() {
hello-20414	3....	1656091.711049:	funcgraph_entry:	0.844 us	copy_strings.isra.26();
hello-20414	3....	1656091.711050:	funcgraph_exit:	1.146 us	}
hello-20414	3....	1656091.711050:	funcgraph_entry:		copy_strings_kernel() {
hello-20414	3....	1656091.711050:	funcgraph_entry:	0.769 us	copy_strings.isra.26();
hello-20414	3....	1656091.711051:	funcgraph_exit:	1.070 us	}
hello-20414	3....	1656091.711051:	funcgraph_entry:		bprm_change_interp() {
hello-20414	3....	1656091.711051:	funcgraph_entry:	0.700 us	kstrdup();
hello-20414	3....	1656091.711052:	funcgraph_exit:	1.107 us	}
hello-20414	3....	1656091.711053:	funcgraph_entry:		open_exec() {
hello-20414	3....	1656091.711053:	funcgraph_entry:	0.397 us	getname_kernel();
hello-20414	3....	1656091.711053:	funcgraph_entry:	4.998 us	do_open_execat();
hello-20414	3....	1656091.711059:	funcgraph_entry:	0.256 us	putname();
hello-20414	3....	1656091.711059:	funcgraph_exit:	6.559 us	}
hello-20414	3....	1656091.711059:	funcgraph_entry:		prepare_binprm() {
hello-20414	3....	1656091.711060:	funcgraph_entry:	0.061 us	mnt_may_suid();
hello-20414	3....	1656091.711060:	funcgraph_entry:	1.196 us	security_bprm_set_creds();
hello-20414	3....	1656091.711061:	funcgraph_entry:	1.799 us	kernel_read();
hello-20414	3....	1656091.711063:	funcgraph_exit:	3.909 us	}
hello-20414	3....	1656091.711064:	funcgraph_entry:		search_binary_handler() {
hello-20414	3....	1656091.711064:	funcgraph_entry:	0.043 us	security_bprm_check();
hello-20414	3....	1656091.711064:	funcgraph_entry:	0.036 us	_raw_read_lock();
hello-20414	3....	1656091.711064:	funcgraph_entry:	0.034 us	try_module_get();
hello-20414	3....	1656091.711065:	funcgraph_entry:	0.050 us	load_script();
hello-20414	3....	1656091.711065:	funcgraph_entry:	0.033 us	_raw_read_lock();
hello-20414	3....	1656091.711065:	funcgraph_entry:	0.035 us	module_put();
hello-20414	3....	1656091.711066:	funcgraph_entry:	0.033 us	try_module_get();
hello-20414	3....	1656091.711066:	funcgraph_entry:	! 248.708 us	load_elf_binary();
hello-20414	3....	1656091.711316:	funcgraph_entry:	0.063 us	_raw_read_lock();
hello-20414	3....	1656091.711316:	funcgraph_entry:	0.036 us	module_put();
hello-20414	3....	1656091.711317:	funcgraph_exit:	! 253.080 us	}
hello-20414	3....	1656091.711317:	funcgraph_exit:	! 271.034 us	}

Easier to understand

```
# trace-cmd report -IS
```

```
CPU 0 is empty
CPU 1 is empty
CPU 2 is empty
cpus=4
```

hello-20414	3....	1656091.711043:	funcgraph_entry:			load_script() {
hello-20414	3....	1656091.711046:	funcgraph_entry:			fput() {
hello-20414	3....	1656091.711046:	funcgraph_entry:	0.143 us		task_work_add();
hello-20414	3....	1656091.711047:	funcgraph_exit:	0.630 us		}
hello-20414	3....	1656091.711047:	funcgraph_entry:			remove_arg_zero() {
hello-20414	3....	1656091.711047:	funcgraph_entry:	0.625 us		get_user_pages_remote();
hello-20414	3....	1656091.711048:	funcgraph_exit:	0.971 us		}
hello-20414	3....	1656091.711048:	funcgraph_entry:			copy_strings_kernel() {
hello-20414	3....	1656091.711049:	funcgraph_entry:	0.844 us		copy_strings.isra.26();
hello-20414	3....	1656091.711050:	funcgraph_exit:	1.146 us		}
hello-20414	3....	1656091.711050:	funcgraph_entry:			copy_strings_kernel() {
hello-20414	3....	1656091.711050:	funcgraph_entry:	0.769 us		copy_strings.isra.26();
hello-20414	3....	1656091.711051:	funcgraph_exit:	1.070 us		}
hello-20414	3....	1656091.711051:	funcgraph_entry:			bprm_change_interp() {
hello-20414	3....	1656091.711051:	funcgraph_entry:	0.700 us		kstrdup();
hello-20414	3....	1656091.711052:	funcgraph_exit:	1.107 us		}
hello-20414	3....	1656091.711053:	funcgraph_entry:			open_exec() {
hello-20414	3....	1656091.711053:	funcgraph_entry:	0.397 us		getname_kernel();
hello-20414	3....	1656091.711053:	funcgraph_entry:	4.998 us		do_open_execat();
hello-20414	3....	1656091.711059:	funcgraph_entry:	0.256 us		putname();
hello-20414	3....	1656091.711059:	funcgraph_exit:	6.559 us		}
hello-20414	3....	1656091.711059:	funcgraph_entry:			prepare_binprm() {
hello-20414	3....	1656091.711060:	funcgraph_entry:	0.061 us		mnt_may_suid();
hello-20414	3....	1656091.711060:	funcgraph_entry:	1.196 us		security_bprm_set_creds();
hello-20414	3....	1656091.711061:	funcgraph_entry:	1.799 us		kernel_read();
hello-20414	3....	1656091.711063:	funcgraph_exit:	3.909 us		}
hello-20414	3....	1656091.711064:	funcgraph_entry:			search_binary_handler() {
hello-20414	3....	1656091.711064:	funcgraph_entry:	0.043 us		security_bprm_check();
hello-20414	3....	1656091.711064:	funcgraph_entry:	0.036 us		_raw_read_lock();
hello-20414	3....	1656091.711064:	funcgraph_entry:	0.034 us		try_module_get();
hello-20414	3....	1656091.711065:	funcgraph_entry:	0.050 us		load_script();
hello-20414	3....	1656091.711065:	funcgraph_entry:	0.033 us		_raw_read_lock();
hello-20414	3....	1656091.711065:	funcgraph_entry:	0.035 us		module_put();
hello-20414	3....	1656091.711066:	funcgraph_entry:	0.033 us		try_module_get();
hello-20414	3....	1656091.711066:	funcgraph_entry:	! 248.708 us		load_elf_binary();
hello-20414	3....	1656091.711316:	funcgraph_entry:	0.063 us		_raw_read_lock();
hello-20414	3....	1656091.711316:	funcgraph_entry:	0.036 us		module_put();
hello-20414	3....	1656091.711317:	funcgraph_exit:	! 253.080 us		}
hello-20414	3....	1656091.711317:	funcgraph_exit:	! 271.034 us		}

Recursive function in the kernel?

The kernel has a fixed stack

Recursive functions can blow that stack

Is this a bug?

Can we exploit it?

What is it doing?

load_script()

```
static int load_script(struct linux_binprm *bprm)
{
    const char *i_arg, *i_name;
    char *cp;
    struct file *file;
    int retval;

    if ((bprm->buf[0] != '#') || (bprm->buf[1] != '!'))
        return -ENOEXEC;
[.]

    /*
     * OK, now restart the process with the interpreter's dentry.
     */
    file = open_exec(i_name);
    if (IS_ERR(file))
        return PTR_ERR(file);

    bprm->file = file;
    retval = prepare_binprm(bprm);
    if (retval < 0)
        return retval;
    return search_binary_handler(bprm);
}
```

Can we exploit this?

```
#!/tmp/blah
```

```
Crash me!
```

```
# echo '#!/tmp/blah  
Crash me!' > /tmp/blah  
  
# chmod +x /tmp/blah  
# /tmp/blah
```

Can we exploit this? Nope!

```
#!/tmp/blah
```

```
Crash me!
```

```
# echo '#!/tmp/blah
```

```
Crash me!' > /tmp/blah
```

```
# chmod +x /tmp/blah
```

```
# /tmp/blah
```

```
bash: /tmp/blah: /tmp/blah: bad interpreter: Too many levels of symbolic links
```


Let's see what happened

```
# trace-cmd record -p function_graph -g 'load_script' sh -c /tmp/blah

  plugin 'function_graph'
sh: 1: /tmp/blah: Too many levels of symbolic links
CPU0 data recorded at offset=0x652000
    0 bytes in size
CPU1 data recorded at offset=0x652000
  155648 bytes in size
CPU2 data recorded at offset=0x678000
   4096 bytes in size
CPU3 data recorded at offset=0x679000
    0 bytes in size
```

Follow the load_scripts

```
# trace-cmd report -O tailprint | grep load_script
```

```
sh-28473 [002] 1660671.155766: funcgraph_entry:      0.122 us | load_script();
sh-28474 [001] 1660671.156664: funcgraph_entry:      | load_script() {
sh-28474 [001] 1660671.156822: funcgraph_entry:      |     load_script() {
sh-28474 [001] 1660671.156999: funcgraph_entry:      |         load_script() {
sh-28474 [001] 1660671.157137: funcgraph_entry:      |             load_script() {
sh-28474 [001] 1660671.157248: funcgraph_entry:      |                 load_script() {
sh-28474 [001] 1660671.157361: funcgraph_entry:      |                     load_script() {
sh-28474 [001] 1660671.157449: funcgraph_exit:      + 88.202 us |                         } /* load_script */
sh-28474 [001] 1660671.157450: funcgraph_exit:      ! 201.306 us |                     } /* load_script */
sh-28474 [001] 1660671.157451: funcgraph_exit:      ! 313.092 us |                 } /* load_script */
sh-28474 [001] 1660671.157451: funcgraph_exit:      ! 452.196 us |             } /* load_script */
sh-28474 [001] 1660671.157452: funcgraph_exit:      ! 630.448 us |         } /* load_script */
sh-28474 [001] 1660671.157453: funcgraph_exit:      ! 770.883 us |     } /* load_script */
```

Follow the load_scripts and a little more

```
# trace-cmd report -O tailprint | grep -C 2 load_script
```

```
CPU 3 is empty  
cpus=4
```

```
sh-28473 [002] 1660671.155766: funcgraph_entry:      0.122 us | load_script();  
sh-28474 [001] 1660671.156664: funcgraph_entry:      | load_script() {  
sh-28474 [001] 1660671.156682: funcgraph_entry:      |     fput() {  
sh-28474 [001] 1660671.156683: funcgraph_entry:      |         task_work_add() {  
--  
sh-28474 [001] 1660671.156821: funcgraph_entry:      0.043 us |  
sh-28474 [001] 1660671.156821: funcgraph_entry:      0.031 us |     _raw_read_lock();  
sh-28474 [001] 1660671.156822: funcgraph_entry:      |     try_module_get();  
sh-28474 [001] 1660671.156822: funcgraph_entry:      |     load_script() {  
sh-28474 [001] 1660671.156822: funcgraph_entry:      |         fput() {  
sh-28474 [001] 1660671.156822: funcgraph_entry:      0.056 us |             task_work_add();  
--  
sh-28474 [001] 1660671.156999: funcgraph_entry:      0.039 us |  
sh-28474 [001] 1660671.156999: funcgraph_entry:      0.032 us |     _raw_read_lock();  
sh-28474 [001] 1660671.156999: funcgraph_entry:      |     try_module_get();  
sh-28474 [001] 1660671.156999: funcgraph_entry:      |     load_script() {  
sh-28474 [001] 1660671.156999: funcgraph_entry:      |         fput() {  
sh-28474 [001] 1660671.156999: funcgraph_entry:      0.048 us |             task_work_add();  
--  
sh-28474 [001] 1660671.157137: funcgraph_entry:      0.033 us |  
sh-28474 [001] 1660671.157137: funcgraph_entry:      0.032 us |     _raw_read_lock();  
sh-28474 [001] 1660671.157137: funcgraph_entry:      |     try_module_get();  
sh-28474 [001] 1660671.157137: funcgraph_entry:      |     load_script() {  
sh-28474 [001] 1660671.157138: funcgraph_entry:      |         fput() {  
sh-28474 [001] 1660671.157138: funcgraph_entry:      0.052 us |             task_work_add();  
--  
sh-28474 [001] 1660671.157248: funcgraph_entry:      0.032 us |  
sh-28474 [001] 1660671.157248: funcgraph_entry:      0.032 us |     _raw_read_lock();  
sh-28474 [001] 1660671.157248: funcgraph_entry:      |     try_module_get();  
sh-28474 [001] 1660671.157248: funcgraph_entry:      |     load_script() {  
sh-28474 [001] 1660671.157249: funcgraph_entry:      |         fput() {  
sh-28474 [001] 1660671.157249: funcgraph_entry:      0.050 us |             task_work_add();  
--  
sh-28474 [001] 1660671.157360: funcgraph_exit:        0.037 us |  
sh-28474 [001] 1660671.157360: funcgraph_entry:      0.032 us |     } /* _raw_read_lock */  
sh-28474 [001] 1660671.157361: funcgraph_entry:      |     try_module_get();  
sh-28474 [001] 1660671.157361: funcgraph_entry:      |     load_script() {  
sh-28474 [001] 1660671.157361: funcgraph_entry:      |         fput() {  
sh-28474 [001] 1660671.157361: funcgraph_entry:      0.040 us |             task_work_add();  
--  
sh-28474 [001] 1660671.157449: funcgraph_exit:        + 15.922 us |  
sh-28474 [001] 1660671.157449: funcgraph_entry:      0.054 us |     } /* prepare_binprm */  
sh-28474 [001] 1660671.157449: funcgraph_exit:        + 88.202 us |     search_binary_handler();  
sh-28474 [001] 1660671.157449: funcgraph_exit:        |     } /* load_script */
```

Follow the load_scripts and a little more

```
# trace-cmd report -O tailprint | grep -C 2 load_script
```

```
CPU 3 is empty  
cpus=4
```

```
sh-28473 [002] 1660671.155766: funcgraph_entry:      0.122 us | load_script();  
sh-28474 [001] 1660671.156664: funcgraph_entry:      | load_script() {  
sh-28474 [001] 1660671.156682: funcgraph_entry:      |     fput() {  
sh-28474 [001] 1660671.156683: funcgraph_entry:      |         task_work_add() {  
--  
sh-28474 [001] 1660671.156821: funcgraph_entry:      0.043 us |     _raw_read_lock();  
sh-28474 [001] 1660671.156821: funcgraph_entry:      0.031 us |     try_module_get();  
sh-28474 [001] 1660671.156822: funcgraph_entry:      |     load_script() {  
sh-28474 [001] 1660671.156822: funcgraph_entry:      |         fput() {  
sh-28474 [001] 1660671.156822: funcgraph_entry:      0.056 us |             task_work_add();  
--  
sh-28474 [001] 1660671.156999: funcgraph_entry:      0.039 us |     _raw_read_lock();  
sh-28474 [001] 1660671.156999: funcgraph_entry:      0.032 us |     try_module_get();  
sh-28474 [001] 1660671.156999: funcgraph_entry:      |     load_script() {  
sh-28474 [001] 1660671.156999: funcgraph_entry:      |         fput() {  
sh-28474 [001] 1660671.156999: funcgraph_entry:      0.048 us |             task_work_add();  
--  
sh-28474 [001] 1660671.157137: funcgraph_entry:      0.033 us |     _raw_read_lock();  
sh-28474 [001] 1660671.157137: funcgraph_entry:      0.032 us |     try_module_get();  
sh-28474 [001] 1660671.157137: funcgraph_entry:      |     load_script() {  
sh-28474 [001] 1660671.157138: funcgraph_entry:      |         fput() {  
sh-28474 [001] 1660671.157138: funcgraph_entry:      0.052 us |             task_work_add();  
--  
sh-28474 [001] 1660671.157248: funcgraph_entry:      0.032 us |     _raw_read_lock();  
sh-28474 [001] 1660671.157248: funcgraph_entry:      0.032 us |     try_module_get();  
sh-28474 [001] 1660671.157248: funcgraph_entry:      |     load_script() {  
sh-28474 [001] 1660671.157249: funcgraph_entry:      |         fput() {  
sh-28474 [001] 1660671.157249: funcgraph_entry:      0.050 us |             task_work_add();  
--  
sh-28474 [001] 1660671.157360: funcgraph_exit:        0.037 us |     } /* _raw_read_lock */  
sh-28474 [001] 1660671.157360: funcgraph_entry:      0.032 us |     try_module_get();  
sh-28474 [001] 1660671.157361: funcgraph_entry:      |     load_script() {  
sh-28474 [001] 1660671.157361: funcgraph_entry:      |         fput() {  
sh-28474 [001] 1660671.157361: funcgraph_entry:      0.040 us |             task_work_add();  
--  
sh-28474 [001] 1660671.157449: funcgraph_exit:      + 15.922 us |     } /* prepare_binprm */  
sh-28474 [001] 1660671.157449: funcgraph_entry:      0.054 us |     search_binary_handler();  
sh-28474 [001] 1660671.157449: funcgraph_exit:      + 88.202 us |     } /* load_script */
```

search_binary_handler()

```
/*  
 * cycle the list of binary formats handler, until one recognizes the image  
 */  
int search_binary_handler(struct linux_binprm *bprm)  
{  
    bool need_retry = IS_ENABLED(CONFIG_MODULES);  
    struct linux_binfmt *fmt;  
    int retval;  
  
    /* This allows 4 levels of binfmt rewrites before failing hard. */  
    if (bprm->recursion_depth > 5)  
        return -ELOOP;
```

I want to know what my computer is doing!

OK, that was interesting (and time consuming)

But what about the rest of my computer?

I want to see it all?

Well, perhaps not all, function tracing is a bit overwhelming!

Just the events please

Let's see the world!

```
# trace-cmd record -e all sleep 10
```

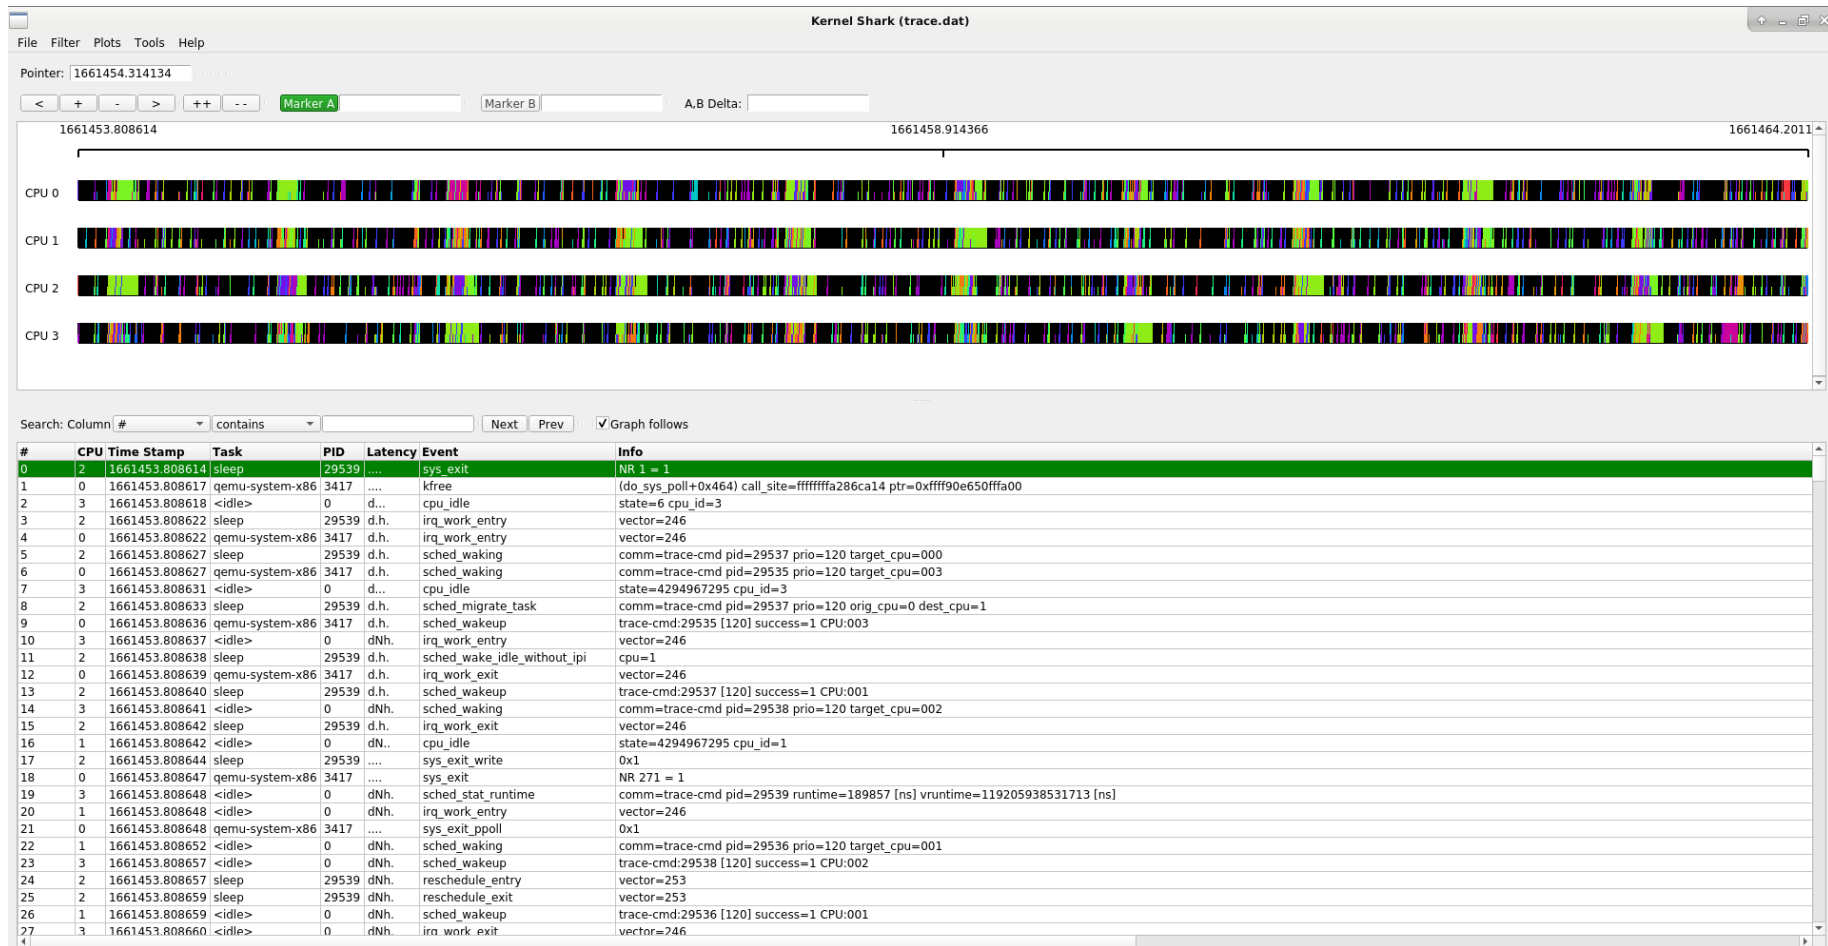
```
CPU0 data recorded at offset=0x822000  
      20647936 bytes in size  
CPU1 data recorded at offset=0x1bd3000  
      21708800 bytes in size  
CPU2 data recorded at offset=0x3087000  
      20733952 bytes in size  
CPU3 data recorded at offset=0x444d000  
      22446080 bytes in size
```

Well, um, that's informative. I think?

trace-cmd report

```
cpus=4
sleep-29539 [002] 1661453.808614: sys_exit: NR 1 = 1
qemu-system-x86-3417 [000] 1661453.808617: kfree: (do_sys_poll+0x464) call_site=ffffffffa286ca14 ptr=0xffff90e650fffa00
<idle>-0 [003] 1661453.808618: cpu_idle: state=6 cpu_id=3
sleep-29539 [002] 1661453.808622: irq_work_entry: vector=246
qemu-system-x86-3417 [000] 1661453.808622: irq_work_entry: vector=246
sleep-29539 [002] 1661453.808627: sched_waking: comm=trace-cmd pid=29537 prio=120 target_cpu=000
qemu-system-x86-3417 [000] 1661453.808627: sched_waking: comm=trace-cmd pid=29535 prio=120 target_cpu=003
<idle>-0 [003] 1661453.808631: cpu_idle: state=4294967295 cpu_id=3
sleep-29539 [002] 1661453.808633: sched_migrate_task: comm=trace-cmd pid=29537 prio=120 orig_cpu=0 dest_cpu=1
qemu-system-x86-3417 [000] 1661453.808636: sched_wakeup: trace-cmd:29535 [120] success=1 CPU:003
<idle>-0 [003] 1661453.808637: irq_work_entry: vector=246
sleep-29539 [002] 1661453.808638: sched_wake_idle_without_ipi: cpu=1
qemu-system-x86-3417 [000] 1661453.808639: irq_work_exit: vector=246
sleep-29539 [002] 1661453.808640: sched_wakeup: trace-cmd:29537 [120] success=1 CPU:001
<idle>-0 [003] 1661453.808641: sched_waking: comm=trace-cmd pid=29538 prio=120 target_cpu=002
sleep-29539 [002] 1661453.808642: irq_work_exit: vector=246
<idle>-0 [001] 1661453.808642: cpu_idle: state=4294967295 cpu_id=1
sleep-29539 [002] 1661453.808644: sys_exit_write: 0x1
qemu-system-x86-3417 [000] 1661453.808647: sys_exit: NR 271 = 1
<idle>-0 [003] 1661453.808648: sched_stat_runtime: comm=trace-cmd pid=29539 runtime=189857 [ns] vruntime=119205938531713 [ns]
<idle>-0 [001] 1661453.808648: irq_work_entry: vector=246
qemu-system-x86-3417 [000] 1661453.808648: sys_exit_ppoll: 0x1
<idle>-0 [001] 1661453.808652: sched_waking: comm=trace-cmd pid=29536 prio=120 target_cpu=001
<idle>-0 [003] 1661453.808657: sched_wakeup: trace-cmd:29538 [120] success=1 CPU:002
sleep-29539 [002] 1661453.808657: reschedule_entry: vector=253
sleep-29539 [002] 1661453.808659: reschedule_exit: vector=253
<idle>-0 [001] 1661453.808659: sched_wakeup: trace-cmd:29536 [120] success=1 CPU:001
<idle>-0 [003] 1661453.808660: irq_work_exit: vector=246
<idle>-0 [001] 1661453.808661: irq_work_exit: vector=246
sleep-29539 [002] 1661453.808661: rcu_utilization: Start context switch
sleep-29539 [002] 1661453.808663: rcu_utilization: End context switch
<idle>-0 [003] 1661453.808664: reschedule_entry: vector=253
<idle>-0 [001] 1661453.808664: rcu_utilization: Start context switch
<idle>-0 [003] 1661453.808666: reschedule_exit: vector=253
<idle>-0 [001] 1661453.808666: rcu_utilization: End context switch
<idle>-0 [003] 1661453.808668: rcu_utilization: Start context switch
sleep-29539 [002] 1661453.808669: sched_switch: trace-cmd:29539 [120] S ==> trace-cmd:29538 [120]
<idle>-0 [003] 1661453.808670: rcu_utilization: End context switch
<idle>-0 [001] 1661453.808670: sched_switch: swapper/1:0 [120] S ==> trace-cmd:29537 [120]
```


What about this?



Introducing KernelShark 1.0

Well, at least it's prettier

Introducing KernelShark 1.0

KernelShark was create in 2009

- GTK version
- Stalled in development
- Slow on large data sets

VMware hired someone to work on it full time

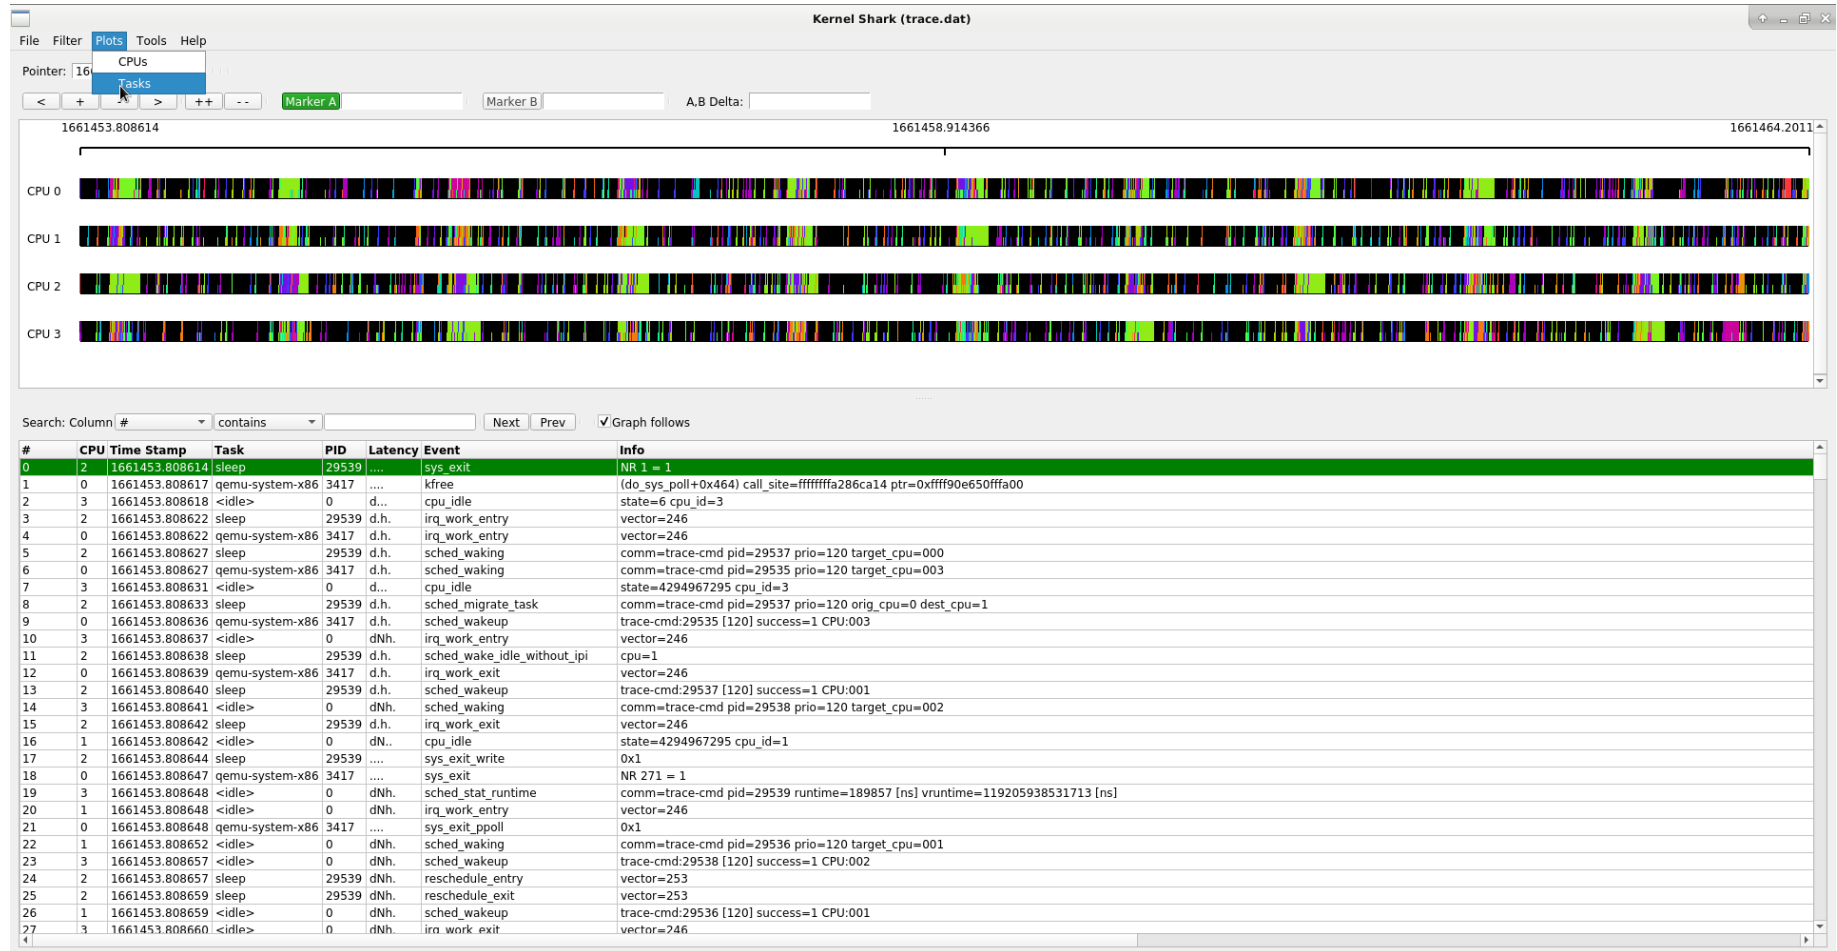
- Rewritten from scratch
- Uses Qt
- Handles large data sets quickly
- Is now a platform (ftrace layer is abstracted out)

Introducing KernelShark 0.99 actually

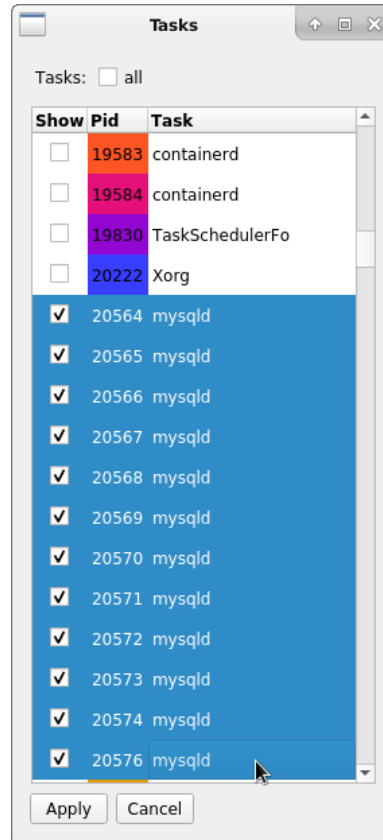
Congratulations!

- You are all now beta testers
- Find a bug, report it here:
<https://bugzilla.kernel.org>
 - Pick “Tools: Tools and utilities”
 - Then “Trace-cmd/Kernelshark”

KernelShark 1.0



KernelShark 1.0 (Task Selection)

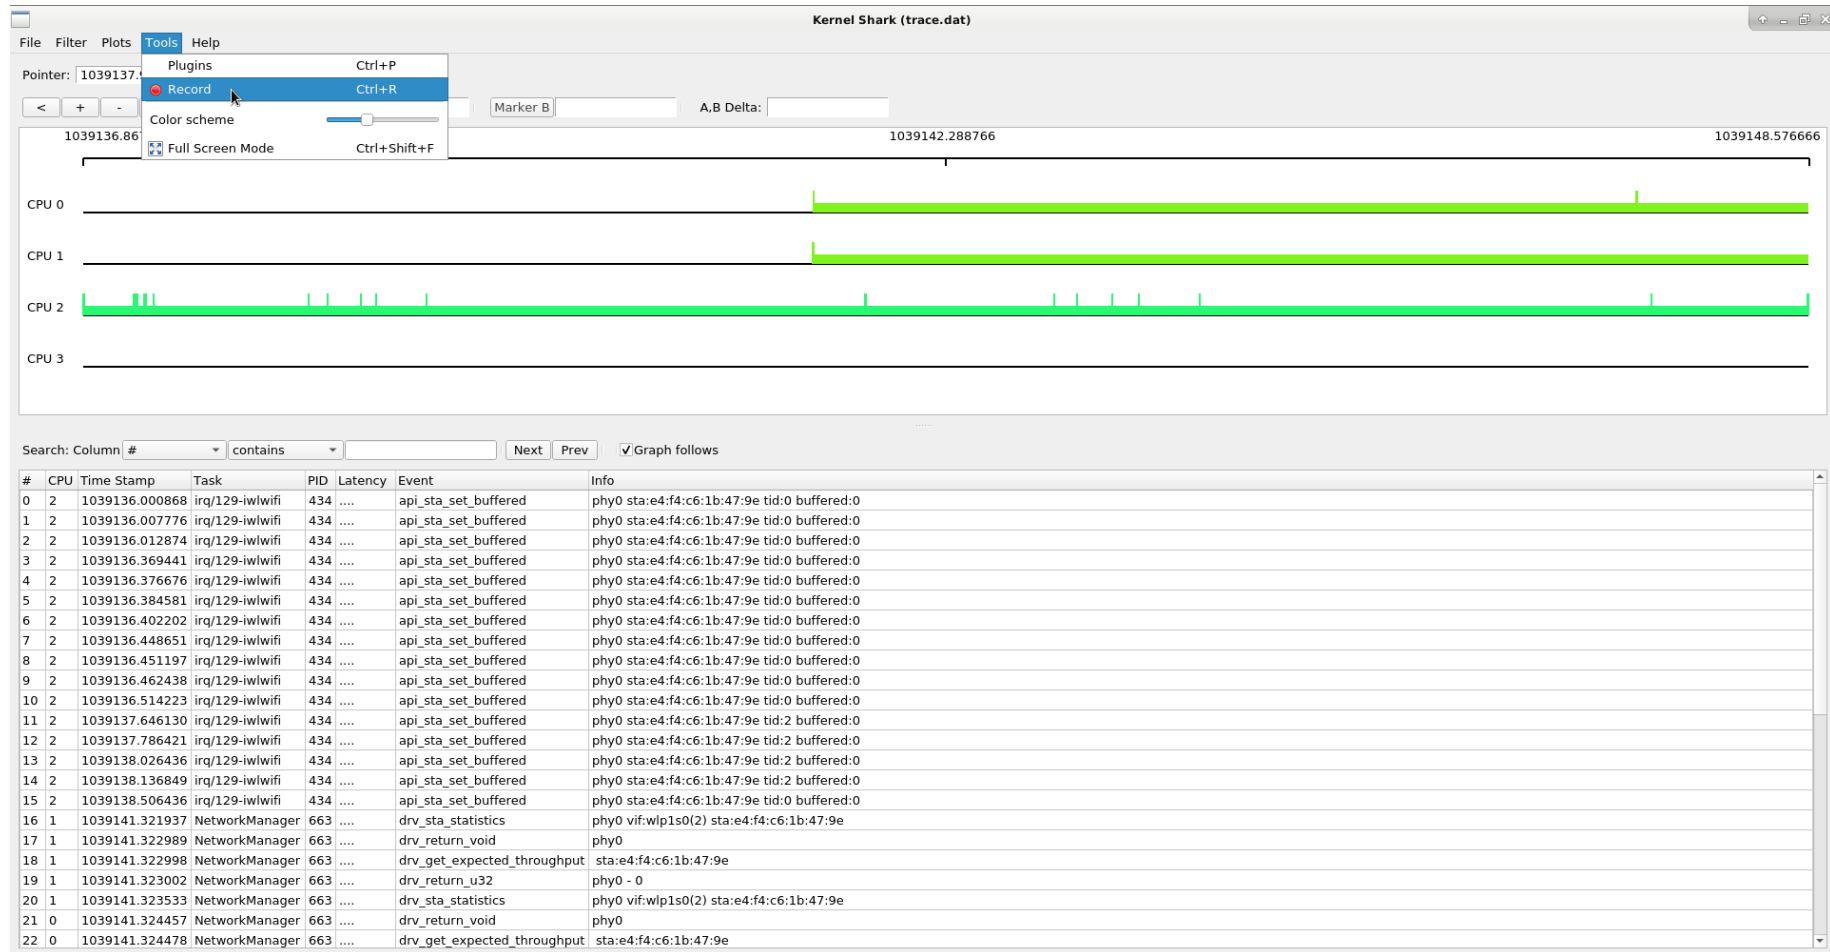


A really simple java program!

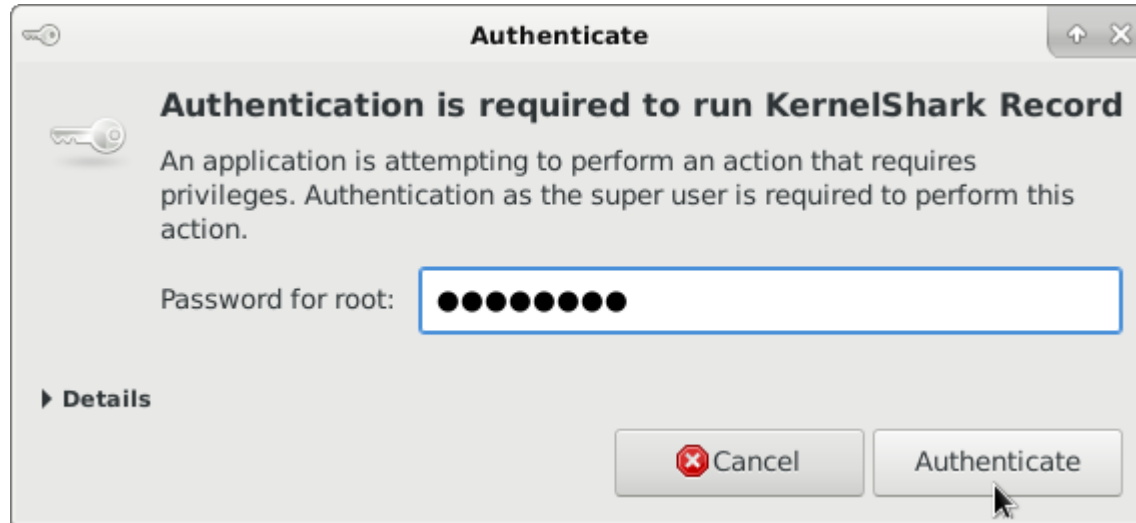
Using your favorite editor, create this file

```
public class HelloWorld {  
  
    public static void main(String[] args) {  
        System.out.println("Hello, World");  
    }  
  
}
```

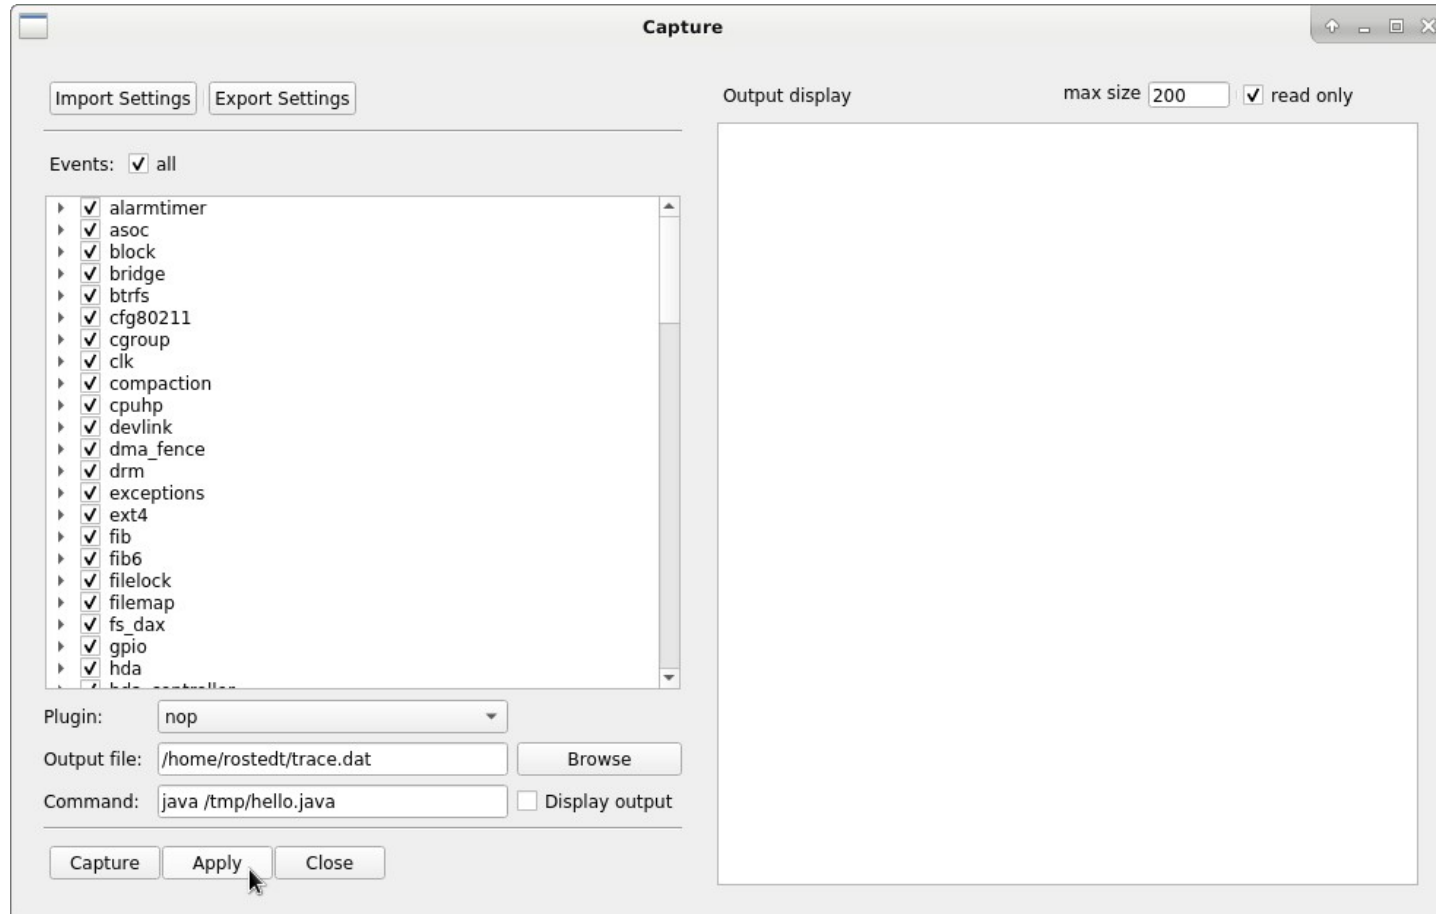
KernelShark 1.0 (Recording)



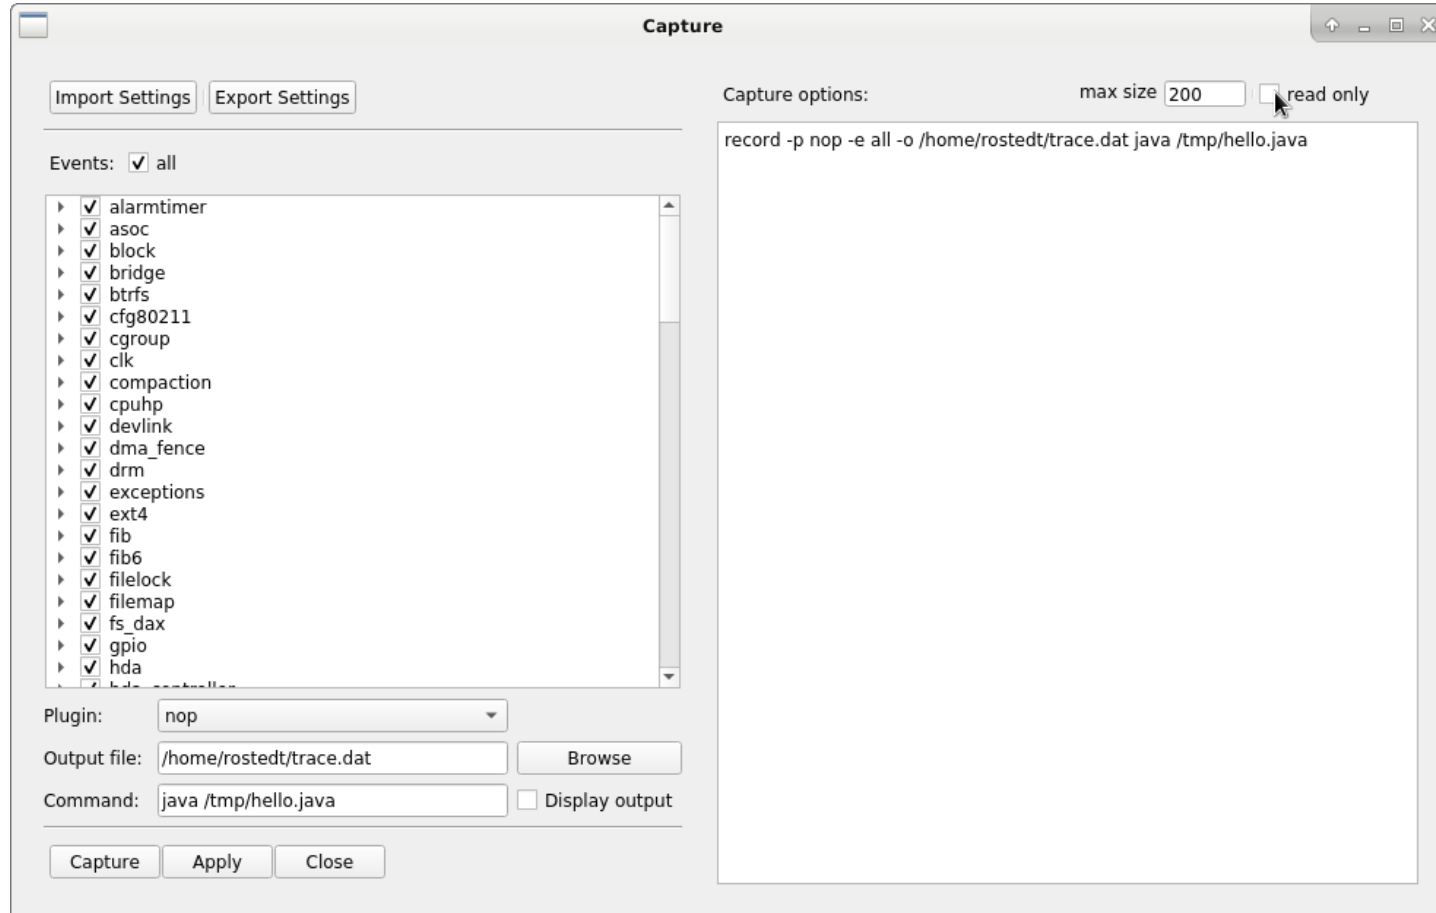
KernelShark 1.0 (Recording)



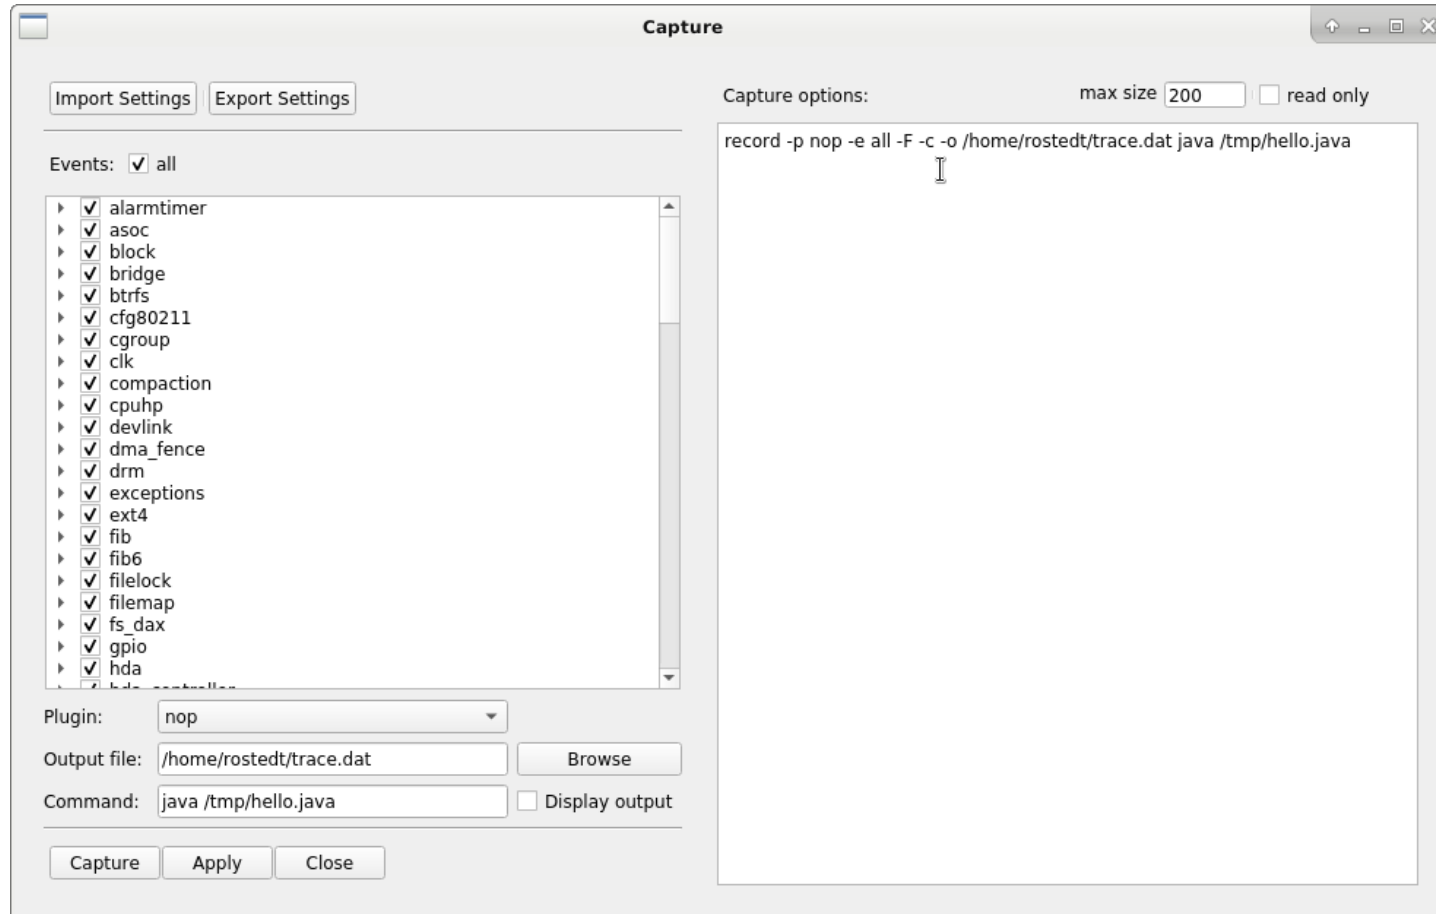
KernelShark 1.0 (Recording)



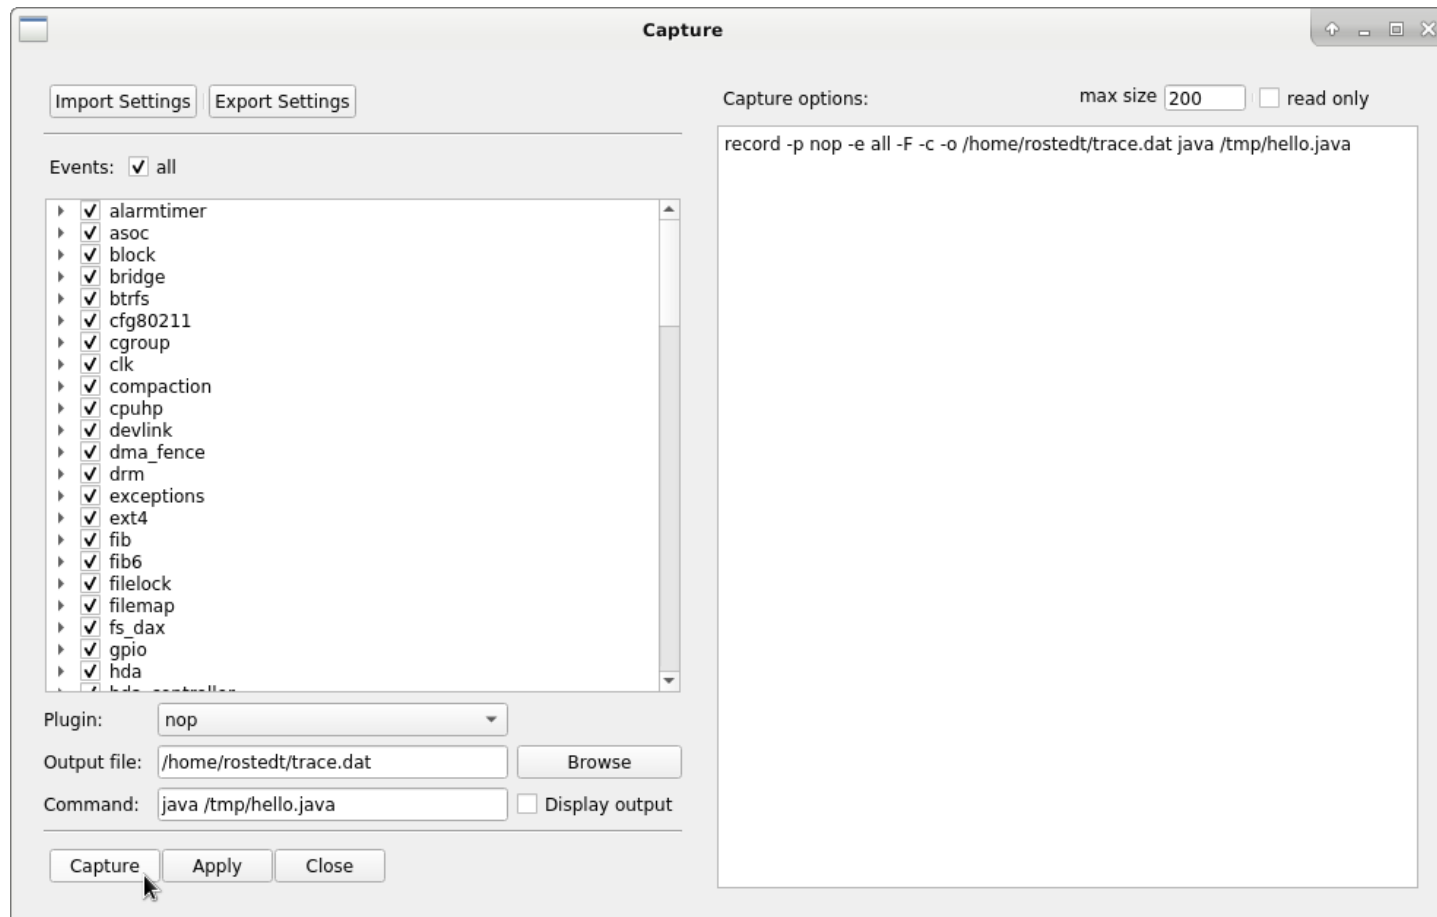
KernelShark 1.0 (Recording)



KernelShark 1.0 (Recording)



KernelShark 1.0 (Recording)



KernelShark 1.0 (Simple Java Program)





Thank You